

OPERATOR'S MANUAL

THE VELOCITY OPERATING SYSTEM

Decision Framework and Field Playbook

The rules, gates, thresholds, run-sheets and templates required to operate a business on the Velocity OS — extracted from the doctrine and rendered as standard work.

Derived from The Velocity Operating System v3.0

Peter V. Griscom, M.Sc.

Chairman, Van Dyke Acquisitions | Founder & Chairman, PVG Capital

FIELD EDITION — FOR INTERNAL DISTRIBUTION

Contents

Sections are coded so they can be cited in meetings and ship lists. "See C.4" is a complete instruction.

Part	Sections
	How to Use This Manual · The System on One Page · The Operating Constitution · The Non-Negotiables · Shared Vocabulary
	A.1 The Decision Sequence · A.2 The Decision Classification Grid · A.3 The Zone Router · A.4 The Convexity Screen · A.5 Survivorship Covenants · A.6 The Decision-Quality Substrate · A.7 The Seven Standing Questions · A.8 The Gate Review Script
	B.1 The Loop at a Glance · B.2 Phase Operator Cards · B.3 The Loop at Three Scales
	C.1 The Stack at a Glance · C.2 Meeting Physics · C.3 Forum Run-Sheets · C.4 Decision Rights
	D.1 Construction Rules · D.2 The Canonical Tree · D.3 Instantiations · D.4 The Tripwire Register · D.5 The Hidden-Economics Briefing · D.6 The Cohort Standard
	E.1 Pipeline Architecture · E.2 Convexity-Adjusted ICE Scoring · E.3 The Pre-Registration Form · E.4 The Six Statistical Non-Negotiables · E.5 Velocity Accounting · E.6 Experimentation Without a Browser
	F.1 The Hiring Loop · F.2 The Outcome Scorecard · F.3 The Candidate Evaluation Matrix · F.4 The Performance Cadence · F.5 The Keeper Test and Separation Protocol · F.6 Development as Capital Allocation
	G.1 Two Rules That Precede Any Scoring · G.2 The Vendor Evaluation Matrix · G.3 Scoring Governance · G.4 Build vs. Buy vs. Orchestrate · G.5 The Agent Org Chart · G.6 The Four Guardrails · G.7 The Revenue-per-FTE Doctrine
	H.1 Panel One — Survivorship · H.2 Panel Two — Throughput · H.3 Panel Three — Velocity & Position
	I.1 The Four Phases · I.2 First Monday Morning
	X.1 The Threshold Register · X.2 The Failure-Mode Index · X.3 Standing Forms · X.4 The Operating Identity

How to Use This Manual

This is not the book. The book argues the case; this manual runs the company. Every page here is an operating instruction: a rule, a gate, a threshold, a run-sheet, a scoring matrix, or a form. Narrative, evidence and case history have been stripped out and left in the source text. What remains is the layer you are expected to execute without re-reading anything.

Read it once end to end. After that, use it the way a pilot uses a checklist — open to the page that governs the decision in front of you, follow it exactly, and close it.

The three tests this manual applies to everything

Test	Question	If the answer fails
Traceability	Does this metric trace through the driver tree to a revenue, margin, opex or cash line in one step?	The metric is deleted. The meeting that reviews it is cancelled.
Zone	Is variance here the waste or the product?	The initiative is misrouted. Re-route at intake before any resource is committed.
Learning yield	Does this increase validated learning per dollar of overhead?	It is waste. It goes.

Reading paths by seat

Your seat	Read first	Then live in
CEO / GM / Owner	Part A (full), Part I (90-day install)	Part C run-sheets; Part H dashboard; the Gate Review Script
Commercial / Sales leadership	Part A §A.2, Part C (Commercial Council, Trade-Spend Gate)	Part D tripwires; Part H Panel Two
Finance / FP&A	Part D (full)	Part D tripwire register; Part H Panels One and Three; Part G TCO discipline
Growth / Marketing	Part E (full)	Pre-registration form; Growth Review run-sheet; Panel Three
People leaders	Part F (full)	Outcome scorecard; candidate matrix; quarterly keeper test
Technology / Ops	Part G (full)	Vendor matrix; agent SLA register; guardrails

Conventions used throughout

Marker	Meaning
	A rule of the system. No exception process exists unless one is written here.
	The priced cost of ignoring the rule. Use this language when you need alignment, not the rule itself.
	A tripwire. When it fires, the default action is stop — not investigate, not discuss.
Gate	A capital-allocation decision with a pre-registered evidence requirement. Default answer is No-

Marker	Meaning
	Go.
Owner	Exactly one name. Never a function, never two people, never a committee.

OPERATING DOCTRINE

Every rule in this manual was written before the situation you are about to apply it to. That is the entire point. A rule renegotiated in the room it constrains is not a rule; it is a preference with a font. If you believe a rule is wrong, change it at the Quarterly Reset, in writing, for everyone — never in the meeting where it is inconvenient.

The System on One Page

One machine. Three engines, one router, one spine, one instrument panel. Everything else in this manual expands one box below.

Layer	What it is	The one-line test
	Waste removal, standard work, error-proofing, survivorship covenants: runway > 12 months, single-position concentration < 25%, fractional Kelly sizing.	Does this remove waste or create ruin risk?
	Barbell. Gaussian zones get Lean. Power-law zones get the largest affordable portfolio of cheap, uncorrelated, capped-downside convex bets. Nothing gets averaged across the two.	Is variance here the waste or the product?
	Intelligence at every workstation. Humans repositioned to judgment, taste, relationships and signature. Headcount decoupled from revenue by design.	What does this workflow become at pennies of intelligence — and who signs it?
	Diagnose → Stabilize & Align → Pilot → Iterate & Scale → Deploy & Lock In. Mandatory loss-framed gates at every phase. Fractal from a one-week test to a multi-market turnaround.	Which phase is this initiative actually in?
	One decomposition of the business. Every scorecard is a horizontal slice of it. Every node has one owner and one owning forum. No residual "other" bucket.	Which P&L line does this number move, and in how many steps?
	Daily protocol and huddle; weekly War Room, Commercial Council, Growth Review, Trade-Spend Gate; monthly MBR; quarterly Reset; annual Offsite. One scorecard, one owner, one decision output each.	Which meeting owns this decision?
	Three panels: Survivorship (are we safe), Throughput (are we learning fast), Velocity & Position (are our bets paying). Leading indicators outnumber lagging two to one.	Which panel moved this week, and why?

Read the economics as a circuit. Lean frees cash from the inside of the business. That cash is the only currency that funds the convex-bet portfolio on the outside. The loop routes every initiative into its correct zone and moves proven winners from the variance zone into the standard-work zone through mandatory gates. The cadence stack is the delivery mechanism; the driver tree guarantees nothing on any scorecard is decorative; the dashboard closes the loop by asking every review the same question in three tenses.

CONSEQUENCE FRAMING

A cadence that generates savings without naming where the freed cash goes is extraction with better hygiene. Kraft Heinz ran the same cost machinery as AB InBev and delivered roughly \$1.7B of savings — then took a \$15.4B goodwill impairment and a 53% stock decline, erasing roughly \$36B of market value, because no forum in the calendar owned the growth question. The reinvestment test is not optional: every forum that frees cash must name the convex destination in the same meeting.

The Operating Constitution

Eight principles, in order of precedence. When two rules in this manual appear to conflict, the higher-precedence principle wins.

#	Principle	What it obligates you to do
1		Model every domain by its actual distribution. Lean the Gaussian zones; barbell the power-law zones. Any plan built on average SKU, average promotion, average account or average pipeline asset is wrong on arrival.
2		No bet, vendor, hire or market entry may create ruin risk. Runway above 12 months, single-position concentration below 25%, fractional Kelly sizing — enforced as covenants, not aspirations.
3		Tests shipped per week, cycle time per loop and cost per experiment lead every lagging financial result by one to three quarters. Manage the lead, not the lag.
4		Every metric on every scorecard traces through the driver tree to a revenue, margin, opex or cash line. Untraceable metrics are deleted.
5		Cohorts over averages. Base rates over anecdotes. Pre-registered criteria over post-hoc stories. The narrative fallacy is a named defect with a named cost.
6		Intelligence at every workstation; humans repositioned to judgment, taste and relationships; headcount growth decoupled from revenue growth by design — and a human gate on anything customer-visible, retailer-visible, financial or legal.
7		Strategy does not cascade through documents. It cascades through a fixed stack of meetings, each with one scorecard, one owner, one decision output.
8		With buyers, partners, subscribers and staff alike, trust leads retention, velocity and compliance. It is built through predictable communication and spent through silence and surprise — which is why it sits on the dashboard next to cash, not in the values statement.

The Non-Negotiables

Rules with no exception process. The moment one becomes negotiable, the system has a hole and entropy will find it.

- One scorecard, one war room, one 7-day ship list. Single-threaded owners on everything.
- No capital, hire or vendor decision without its matrix, its pre-mortem and its decision-journal entry.
- Cohorts over averages; base rates over anecdotes; pre-registration over post-hoc narrative.
- Runway > 12 months; single-position concentration < 25%; fractional Kelly sizing; no ruin, ever.
- Pilot before scale — for channels, products, promotions, vendors and people alike.
- Kill fast, publicly and cheaply. Celebrate the kill as validated learning.
- Ship the daily structural-hole contact and the reputation asset before reactive work begins.
- One play at a time through the Commercial Council. A business running five plays is running none.
- Every AI workflow has a named human owner and a stated gate before it ships — not after the first complaint.

OPERATING DOCTRINE

An initiative that cannot name its owner, its WIP slot and its kill criterion does not exist. Neither does a meeting

that cannot name its decision output.

Shared Vocabulary

The loop is fractal, which means a coordinator at any level can walk into any review and ask the same questions. That only works if the words mean the same thing in every room. These definitions are binding.

Term	Binding definition
Binding constraint	The single variable whose relief moves the P&L most. Named on one page with cohort evidence, not five priorities confirmed with enthusiasm.
Gate	A capital-allocation decision with a pre-registered evidence requirement. The default answer is No-Go; the burden of proof sits with continuation.
Pre-registration	Hypothesis, single target metric, minimum detectable effect, sample size, run time, guardrails, kill criteria and the decision each outcome triggers — all written and signed before launch.
Convexity	Bounded, known downside against structurally unbounded upside. A position with capped upside and open-ended downside is a position to exit, not to fix.
Zone	Inside the machine (variance is waste, apply Lean) or outside the machine (variance is the product, apply convexity). Decided once, at intake.
Ship list	The binary seven-day commitment set produced by the War Room. Shipped or not shipped. No partial credit, no status colours.
Play	The single route-to-market execution the commercial organisation runs this week, with adoption tracked by account and by rep.
Tripwire	A pre-set threshold whose breach changes the default action without further debate. Two tripwires firing in one quarter triggers a mandatory standing question.
Single-threaded owner	Exactly one accountable name per lever, experiment, metric and initiative. Shared ownership is unowned.
Standard work	A validated win codified, documented, trained and running without its originator present. A win that never becomes standard work was a win that never happened.
Decision journal	The permanent record of material decisions: predicted outcome, evidence, matrix scores, pre-mortem, and the name that signed. Renewals and retrospectives are judged against recorded predictions, never reconstructed memory.
Type 1 decision	High-impact and effectively irreversible. Requires the full substrate: pre-mortem, red team, checklist, journal entry.
Type 2 decision	Reversible at bounded cost. Delegate to the single-threaded owner and move; the cost of deliberation exceeds the cost of being wrong.
Reinvestment test	Every forum that frees cash names where the cash goes in the same meeting. If the destination is not the convex portfolio, the cadence is extraction.

PART A

The Decision Framework

Part A is the router. Before any initiative consumes a dollar, an hour or a seat, it passes through five screens in fixed order: classification, zone, convexity, survivorship, and substrate. The screens are ordered so that the cheapest disqualifier fires first. Most proposals die on screen one or two, which is the intended yield.

A.1 The Decision Sequence

Screen	Question	Disqualifier	Governing artifact
1. Classify	Is this Type 1 (irreversible, material) or Type 2 (reversible, bounded)?	Type 2 decisions escalated upward — that is queue-building, and queues are waiting waste.	Decision Classification Grid (A.2)
2. Route	Is variance here the waste or the product?	Venture logic applied to operations, or Lean logic applied to the bet portfolio.	Zone Routing Table (A.3)
3. Price the shape	Is the payoff convex, symmetric or concave?	Capped upside with open-ended downside. Exit, do not optimise.	Convexity Screen (A.4)
4. Test survival	If this fails completely, can we survive it? Does it breach a covenant?	Any breach of runway, concentration or Kelly limits. No override authority exists.	Survivorship Covenants (A.5)
5. Apply substrate	What evidence, dissent and documentation does this class of decision require?	Material commitment with no pre-mortem, no red team and no journal entry.	Decision-Quality Substrate (A.6)

OPERATING DOCTRINE

Screens are run in order and the first failure ends the process. An initiative that fails screen 4 is not improved by a brilliant answer to screen 5. The most expensive meetings in any company are the ones spent perfecting proposals that should have died on screen two.

A.2 The Decision Classification Grid

Speed is impossible without pre-cleared authority, because every uncleared decision right is a queue with a measurable opex cost. Classify first; the class determines who decides, how fast, and what paperwork attaches.

Class	Test	Who decides	Required artifacts	Latency standard
	Irreversible or reversible only at severe cost; > 10% of quarterly discretionary spend; touches a covenant; customer-, retailer-, regulator- or capital-facing.	Named executive with formal red-team dissent on the record; CEO signs.	Matrix score, pre-mortem, red-team memo, base-rate check, decision-journal entry.	≤ 10 business days from proposal to signature.
	Changes the operating system itself: cadence, decision rights, covenants,	CEO at the Quarterly Reset only. Never mid-	Written proposal, impact on each of the eight principles,	Quarterly cycle. No emergency

Class	Test	Who decides	Required artifacts	Latency standard
	org design, zone routing rules.	quarter, never in the meeting it constrains.	transition plan.	path.
	Reversible at bounded, known cost. Fits inside an existing gate, budget and WIP slot.	The single-threaded owner. Alone. No committee.	Ship-list entry. Journal entry only if the outcome surprised the owner.	≤ 3 business days. Silence from Consulted parties is consent.
	A bounded purchase of information: pilot, test market, paid vendor trial, work sample.	Growth owner or initiative owner, funded by conviction.	Pre-registration doc signed before launch. Nothing else.	Same week. Conviction is sufficient budget authority at pilot scale.
	A tripwire fired, a gate was missed, or a covenant is within tolerance of breach.	Nobody. The hold is automatic; only the named covenant owner may release it.	Release memo naming the condition that changed.	Immediate. The hold precedes the discussion.

CONSEQUENCE FRAMING

In stalled businesses the ten most recent material decisions typically required 6–9 touches and 30–70 elapsed days. After decision rights are published with latency standards, the same decision classes run at 2–3 touches and 3–10 days. That delta is not administrative housekeeping — revenue delayed is revenue destroyed at the discount rate, and every uncleared authority is a permanent queue billed to opex every week it exists.

A.3 The Zone Router

The routing question is asked once, at intake, and the answer determines everything downstream: the discipline applied, the scorecard, the cadence and the gate criteria. Get this wrong and operational excellence will simply make you fail more efficiently.

Initiative type	Zone	Governing discipline	Gate
Promotion mechanics, discount depth, trade calendar	Inside	Lean: event-level ROI scorecard, error-proofed accruals	Event ROI ≥ threshold at day 60–90 net of decay; kill or redesign losers at the monthly review
Supply chain, S&OP, fulfilment, service delivery	Inside	Lean: standard work, OEE, forecast-error reduction	Forecast error < 10%; performance sustained two review cycles
Compliance, claims, financial control, clinical or credit process	Inside	Lean: error-proofing, stop-the-line authority, zero-defect standard	Zero open severity-1 items; audit sample clean
Pricing and pack architecture tests	Outside → in	Convexity: bounded A/B tests, pre-registered MDE	Statistically meaningful movement in 30 days; winner locks as list-price standard at Phase 5
New brand, category, market or product bets	Outside	Convexity: fractional Kelly sizing, staged test-market sequence	Velocity wins proven in test markets before distribution dollars; kill criteria pre-registered

Initiative type	Zone	Governing discipline	Gate
Content and authority engine	Outside	Convexity: cheap shots, high volume, power-law reach expectations	Cost per asset falling; reach outliers reinvested, median outputs killed
M&A, partnerships, capital deployment	Outside, with survivorship covenants	Convexity plus concentration caps: earnout structures, downside capped by structure	No single position > 25% of enterprise value; downside capped at structure, never at hope

The two failure directions

Direction	What it looks like	Timeline to damage	Countermeasure
	Timid, over-analysed, consensus-driven innovation. Every concept must prove itself to an average before funding; every test is resized until it is safe enough to be meaningless. The portfolio converges to line extensions and promotion overlays and the growth algorithm quietly becomes price-over-volume.	Slow — 3 to 5 years, arriving as a stall with margins at multi-year lows.	Kill-rate floor of 50% in the pipeline. If the variance zone has produced no kills in two quarters, you are not running enough shots.
	Undocumented commitments, unmeasured spend, tail products launched on enthusiasm, cash tied up in invalidated inventory. Every promotion is "strategic" and none has an accrual matched to it.	Fast — the cash crisis arrives on schedule and forces premature exits from positions that only needed time.	Event-level ROI on every discretionary spend; the Trade-Spend Gate; weekly covenant review.
	The null case: no operating system compressing cycle time on the inside, no gated, kill-capable discipline on the capital-allocation side. Buybacks, acquisitions and reserves substitute for a machine.	Very slow, then all at once — the largest recorded value destructions in modern industry sit here.	Install both engines in the same 90 days. Neither substitutes for the other.

A.4 The Convexity Screen

Score every proposed position before it enters the backlog. The screen is deliberately crude — its job is to sort, not to price.

Dimension	Score 1	Score 3	Score 5
Downside bound	Open-ended; failure can impair the balance sheet or the brand	Bounded but large; failure costs a quarter	Capped at a pre-agreed pilot cost known before launch
Upside structure	Fixed, one-time step change	Repeatable but linear in spend	Structurally unbounded; wins compound on every future cohort or order
Reversibility	Multi-year lock-in; unwinding costs more than the entry	Reversible within a year at moderate cost	Reversible inside one cycle at near-zero cost
Information yield	Outcome will be unreadable regardless of result	Result readable but confounded	Answers a named driver-tree question for a budgeted price
Correlation	Correlated with existing	Partially correlated	Uncorrelated with the existing

Dimension	Score 1	Score 3	Score 5
	exposures; fails when everything else fails		book

- Composite above 3.0 is required for any initiative entering the outside zone. Below 3.0 the position is symmetric or concave and belongs inside the machine, if it belongs anywhere.
- Portfolio standard: more than 60% of the active test pipeline must be structurally unbounded. Below that, the pipeline is polishing rather than positioning.
- A convexity score of 5 on downside bound with a 1 on information yield is the most common trap — a cheap test that cannot teach you anything is not cheap, it is free waste.

OPERATING DOCTRINE

Conviction may fund the pilot. Only data may fund the scale-up. The pilot is cheap, fast and reversible, so conviction is sufficient budget authority. The scale-up is expensive, slow to reverse and organisationally self-sealing, so only measured lift against a control earns it. Operators who let conviction fund scale-ups are gambling; operators who demand full significance for pilots are frozen. The gate is the exchange rate between the two currencies.

A.5 Survivorship Covenants

These are balance-sheet covenants, not targets. They are reviewed weekly, attested monthly by name, and have no override authority. A covenant with an override is a guideline.

Covenant	Limit	Review	Action on breach
Cash runway (13-week rolling)	> 12 months	Weekly review, monthly attestation	Auto-hold on all outside-zone deployment. Reset the bet portfolio before anything else is discussed.
Single customer / channel concentration	< 25% of net revenue; flag at 20%	Monthly attestation	Named diversification plan with dates at the next MBR. No new concentration-increasing commitments.
Single position vs. enterprise value	≤ 25%	Quarterly audit	Structure down: earnouts, staged tranches, partner participation. Do not proceed on hope.
Position sizing	All bets ≤ fractional Kelly cap	Quarterly audit	Resize before the next tranche deploys.
Vendor / model concentration for critical workflows	No single provider > 25% single point of failure	Quarterly	Stand up a second path before renewal, not after the outage.
Compliance and quality incidents	Zero open severity-1	Weekly tripwire	Stop-the-line. Nothing ships in the affected workflow until closed.

CONSEQUENCE FRAMING

Concentration is the covenant that converts slowly everywhere else and instantly at the end. A single-counterparty limit of 25% of capital is federal banking law for exactly this reason. Large suppliers operate with a single retailer at

roughly 16–20% of net sales as a deliberately managed dependency — survivable at their scale, fatal at yours, because one planogram reset becomes a company-level event. Manage the number, or discover it in a crisis.

A.6 The Decision-Quality Substrate

Both zones fail cognitively before they fail financially. The Lean zone fails by normalising drift — the forecast miss that becomes "seasonality." The variance zone fails by narrative capture — the sunk pilot that becomes "strategic." Four artifacts are the only durable countermeasure, and they are gating infrastructure, not culture-building suggestions.

Artifact	When mandatory	Standard of completion
	Every material commitment, without exception.	"It is twelve months from now and this failed. Why?" The three most-cited failure modes each get a named owner and a leading metric designed to surface them early. Filed before signature.
	Every Type 1 decision.	A formally resourced dissenting case, written by someone whose job that week is to defeat the proposal. Not a devil's advocate at the table — a memo, in advance, on the record.
	Every gate.	Bias check, base-rate check, consequence framing, reversibility classification, and an explicit answer to "what have you done to prove your thesis wrong?"
	Every Type 1 decision and every surprise.	Predicted outcome with a date, evidence relied on, matrix scores, the pre-mortem, and the name that signed. Reviewed against reality at renewal or retrospective.

The base-rate check

The cheapest discipline available and the one most often skipped. Before approving anything, ask: what happened the last ten times a company like ours tried this? It costs one meeting and prevents the most expensive class of error. Where an outside view exists, it outranks the inside view by default; the burden of proof sits with the person claiming this time is different.

A.7 The Seven Standing Questions

Asked at every gate, every business review, and every decision that moves capital, people or positions. Not a checklist to complete once — the permanent test pattern the operating system runs against reality.

#	Question	What it catches
1		Hidden economics. If the spike decays into a lower baseline — deal-dependent customers, trained partners, diluted price integrity — the volume was borrowed, not earned.
2		Capped upside with open-ended downside: the signature of a position to exit, not to fix.
3		An experiment without a pre-registered criterion is not an experiment; it is a narrative waiting to be written.
4		Ruin risk. No bet, hire, vendor or market entry may create it. The

#	Question	What it catches
		pre-mortem is mandatory insurance.
5		Confirmation is the default cognitive setting; disconfirmation is the work. If nothing has been done to falsify the thesis, the thesis is still a story.
6		The keeper test, run in both directions. A "maybe" is a "no" with a delay attached.
7		Re-architecture candidates, and the guardrail that keeps leverage from becoming liability.

A.8 The Gate Review Script

Every gate in this manual is run with the same six moves, in the same order, by whoever chairs the forum. The script exists so that gate discipline does not depend on the seniority or mood of the person holding it.

Move	Words	Why
1. Read the criterion aloud	"The criterion signed on [date] was: [exact text]."	Pre-registration only works if it is read from the document, not from memory. Reading it aloud removes the option of a quiet redefinition.
2. State the result	"The measured result is: [number]. Against a required [number]."	Result before interpretation. Numbers first, always.
3. Call the verdict	"That is a kill / iterate / scale."	Three verdicts exist. "Keep watching" is not one of them — it is how zombies consume budget.
4. Ask the override question	"What would have to be true for us to override our own criterion?"	The only legitimate route to continuation. If nobody can answer it in one sentence, the verdict stands.
5. Name the next commitment	"Owner, resource, and date for the next loop."	A gate that ends without a named next commitment has produced awareness, not a decision.
6. File it	"Journal entry and prediction recorded."	Converts the gate from an event into longitudinal data. Renewal decisions are made against recorded predictions, never reconstructed memory.

FROM THE FIELD

The meeting that changes an organisation's reflexes is almost never the one where something is approved. It is the first one where a pre-registered floor is enforced against a senior voice asking for one more quarter. Once the room watches a gate hold under pressure, the gates become real everywhere else at zero additional cost — and the loudest advocate for one more quarter frequently becomes the most aggressive killer of weak initiatives in the building.

PART B

The Execution Loop

Every initiative in the enterprise runs through one five-phase loop with one set of gates. A brand turnaround, a product launch, a vendor implementation, a hiring process, a single week's checkout experiment — same loop, same gate questions, only the nouns change. Fractality is what allows a lean executive team to run dozens of concurrent loops on one shared mental model.

The loop is not elegance. It is economics: businesses die from a small set of recurrent diseases — portfolio sprawl, promotion addiction, distribution expanded ahead of proven velocity — and all of them are the same underlying defect wearing different labels: **activity that consumes cash without compounding the repeat engine**.

B.1 The Loop at a Glance

Phase	Objective	Core artifacts	Gate criteria (Go / No-Go)
	Find the binding constraint in ≤ 10 days using cohort data, not averages. Run the trust and confidence scan across partners, key accounts and the front line.	Scorecard v1; segmentation by cohort and account; rumour map; hidden-economics check	Binding constraint confirmed with cohort data; scorecard live with ≤ 7 metrics, each with owner and baseline; decision rights published
	Stop value erosion. Install the operating spine. Concentrate the organisation on the vital few.	War room stood up; spend gate live; stop-doing list issued; core offer architecture defined; commercial compact signed	Cash and trust bleed demonstrably slowed; WIP limit enforced; single-threaded owners named
	Test the intervention on a bounded cohort against pre-registered success criteria before spending scale dollars.	Experiment docs with hypothesis, MDE and kill criteria; pilot cohort definition; matched controls or baseline period	Statistically meaningful movement in the target leading metric within 30 days, measured against control
	Expand what worked, kill what did not, and re-validate economics at every expansion ring.	Weekly ship list; ring-by-ring scaling plan; incentive-overlay economics validation	Unit economics hold at each ring; overlay / incentive ROI validated $\geq 1.5\times$; no unintended-consequence flags
	Convert validated wins into standard work, codified playbooks and permanent governance.	Codified operating model; sunset-or-formalise plan for temporary structures; continuous-improvement calendar	Playbooks documented and trained; gains sustained ≥ 2 review cycles without founder intervention

OPERATING DOCTRINE

Read the gate column as the operating system and the rest of the table as documentation. Each gate converts a phase from an activity into a capital-allocation decision with an explicit evidence requirement. The default answer at every gate is No-Go, and the burden of proof sits with continuation — which inverts the sociology of most organisations, where launches are celebrated and kills are career events.

B.2 Phase Operator Cards

One card per phase. Each card is written to be executed by someone who has not read the source book.

R — Diagnose | Ten days to the binding constraint

Field	Standard
Time box	10 days. Non-negotiable. Gate convenes on day 10 regardless of how interesting the remaining questions are.
Owner	Transformation lead at ~80% allocation, supported by one FP&A analyst at ~50% and the commercial lead for the trust scan.
Method	Pull cohort data, never averages: repeat behaviour by acquisition cohort, performance by account and by unit, contribution by product and by channel. Reconstruct the last ten material decisions from calendar and email to measure touches and elapsed days.
Hidden-economics check	Decompose apparent lift into subsidisation of buyers who would have paid full price, cannibalisation of your own volume, forward-buying, and genuine incrementality. Ask the standing question in writing: are we buying today's volume with tomorrow's baseline erosion?
Trust scan	Short anonymous instrument across the front line and the top accounts. One question does most of the work: were commitments this company made in the last four quarters kept?
Output	One page. The binding constraint, named, with the cohort cut that proves it. Scorecard v1 live with ≤ 7 metrics. RACI published.
	The diagnostic expands to fill available curiosity. Week three arrives with a beautiful segmentation and no decision.
	The ten-day clock. Anything not bearing on the constraint goes to a parking lot that nobody is allowed to work.

A — Stabilize & Align | Stop the bleed first

Field	Standard
Time box	4–6 weeks.
Rule	You do not grow your way out of a leak; you plug it. No growth dollar is spent until the bleed has demonstrably slowed.
Move 1 — Spend gate	Every discretionary commitment above the materiality threshold requires a 60–90 day true-ROI projection — lift minus subsidisation, cannibalisation and post-event decay — signed jointly by finance and the business owner.
Move 2 — Stop-doing list	Publish it. Tail products, zombie programs, meetings with no decision output. Stopping is the fastest capacity you will ever create, and the tail does not merely underperform — it corrupts every average the organisation manages to.
Move 3 — Commercial compact	One play per week with participation tracked, replacing the previous pattern of concurrent "priorities." WIP limit of five initiatives with an intake gate: new work enters only by displacing something on the stop-doing list.
Move 4 — Predictable	Trust is rebuilt by predictable communication and spent by silence. Publish the cadence to partners and the

Field	Standard
communication	front line and then never miss it.
Gate	Cash bleed quantified and slowed; WIP limit enforced; single-threaded owners named on every surviving initiative.
	Skipping straight to growth. Pilot results measured against a still-bleeding baseline are unreadable, and a front line asked to execute growth plays while firefighting does neither.
	Stabilize is not caution. It is sequencing. Gate 2 does not pass on effort.

P — Pilot | Bounded cohort, pre-registered criteria

Field	Standard
Time box	30 days. No extensions, no "one more month," no reframed metrics.
Owner	Growth owner at ~60% allocation; transformation lead shifts to cadence enforcement.
Scope	A bounded cohort with matched controls. Two experiments against the constraint, not twelve.
Pre-registration	Written and signed before launch: falsifiable hypothesis, the single driver-tree metric it must move, minimum detectable effect, required sample and run time, guardrail metrics, kill criteria, and the decision each outcome triggers.
Funding rule	Conviction is sufficient budget authority at pilot scale. It is never sufficient at scale-up.
Gate	Movement at or beyond the registered MDE against control, inside the window. A miss is a kill or a re-pilot — never a rename.
	Criteria drift. The pilot misses its bar and the room renegotiates the definition of success after the fact.
	The experiment doc signed on day one is read aloud at the gate, verbatim, before any result is discussed.

I — Iterate & Scale | Ring by ring, economics at each step

Field	Standard
Method	Expand in rings — by banner, region, channel or segment — because economics degrade as you leave your best-fit customers. Each ring must re-earn its expansion.
Three questions per ring	Does the leading metric hold? Do unit economics hold? Does the incentive or overlay ROI stay above 1.5×?
Kill discipline	A ring that misses its floor is killed inside three weeks, freeing working capital rather than stranding it. Kill publicly and cheaply; the kill announcement is culture-building, not embarrassment.
Sequencing law	Velocity before distribution. Prove per-unit productivity in the current ring before buying the next tranche of reach. Distribution can be purchased; velocity must be earned, and it compounds — a buyer watching outlier per-door productivity hands you the next tranche on favourable terms.
Gate	Unit economics hold at each ring; overlay ROI $\geq 1.5\times$ on attributable margin or the overlay is redesigned or terminated; governance adopted — the weekly scorecard meeting has run four consecutive weeks with decisions logged.

Field	Standard
	Premature scale. Ring one is launched on pilot enthusiasm rather than pilot evidence, and a bad result in a major account costs trust that took eight weeks to rebuild.
	Ring gating. No ring opens until the prior ring's economics are on the scorecard, and the CEO personally delivers the "not yet" to the account asking for expansion.

D — Deploy & Lock In | Standard work or it didn't happen

Field	Standard
Objective	Convert wins into institutional property. The transformation lead returns to point, now deliberately working themselves out of a job.
Codify	Every validated play becomes standard work: written, trained, and in use by people who were not in the pilot.
Sunset	All temporary structures are formalised or closed. The war room either becomes the standing operating meeting or it ends.
Re-aim	Run the Quarterly Reset: convexity audit of the full pipeline, threshold-progress review, and capital reallocation toward what the gates validated. Lock-in is not freezing the machine; it is wiring the reset loop so the machine keeps re-aiming itself.
Gate	The operating model runs two consecutive full cycles — scorecard meeting, ship list, gate review — without founder intervention, with gains sustained inside tolerance.
	Founder gravity. The CEO keeps re-entering decisions the system is designed to make, and the organisation correctly concludes the cadence is theatre.
	The taper is contractual. The CEO's gate role shifts from deciding to auditing decision quality, and any override is logged and reviewed at the next reset.

OPERATING DOCTRINE

An initiative that cannot show pre-registered metric movement at the Pilot gate is killed or restructured. Continuing it is reclassified from "perseverance" to a new, and stupid, capital-allocation decision. Sunk costs are recovered exactly never, and every week a failing program consumes is a week of compounding donated to the competitor who killed their version of it faster.

B.3 The Loop at Three Scales

Same five phases, same gates, three orders of magnitude apart. Use this table to answer the most common objection — "our thing is different."

Phase	Multi-market turnaround	Vendor implementation (~\$50K)	One-week growth experiment
Diagnose	Cohort decomposition across accounts, products and channels; trust scan; hidden-economics check.	Which workflows actually bleed margin or cycle time today?	Which driver-tree node is underperforming, and by how much?
Stabilize	Spend gate live; stop-doing list; WIP limit;	Stop the spreadsheet shadow system; freeze process changes.	Freeze concurrent changes to the same surface; establish clean

Phase	Multi-market turnaround	Vendor implementation (~\$50K)	One-week growth experiment
	commercial compact.		baseline.
Pilot	Two banners or segments, 140 matched controls, 30 days.	Two divisions, pre-registered adoption criteria, 30–60 days, paid.	Bounded audience, MDE and kill criteria set in advance.
Iterate	Ring by ring with economics re-validated at each.	Expand by region, re-validating adoption at each step.	Iterate the variant or kill it. No third option.
Deploy	Playbooks, standard work, war room sunsetted.	Codified into the monthly business review.	Winner becomes standard; result filed in the failure archive either way.

CONSEQUENCE FRAMING

Speed is the strategic variable, not the plan. Revenue delayed is revenue destroyed at the discount rate; trust in partners and the front line decays measurably during leadership silence; and every week of deliberation is a week of compounding donated to competitors. Speed is also what makes honest experimentation base rates survivable: if roughly one in three tests wins, the team running twelve loops a quarter finds four winners while the annual-planning team is still wordsmithing its first hypothesis.

PART C

The Cadence Stack

Strategy does not cascade through documents. It cascades through a fixed architecture of meetings, each with one scorecard, one owner, one decision output and a hard time box. No stalled business has ever recovered on the strength of a strategy deck. Many have recovered on the strength of a calendar.

The stack is deliberately boring: same rooms, same scorecards, same decision outputs, every week, until the organisation runs it without thinking. Boring is the design goal.

C.1 The Stack at a Glance

Cadence	Forum	Duration	Owner	Decision output
Daily	Morning Protocol (individual): survivorship check, convexity prioritisation, one bridge-building contact before reactive work	45 min	Each executive	Today's single highest-convexity action, time-blocked
Daily	Team huddle / stop-the-line review: yesterday's ship list, today's blockers, any stop-the-line flags	15 min	Single-threaded owners	Unblocked ship list; escalations routed
Weekly	Executive War Room: scorecard movement → next 7-day ship list	30–45 min	CEO/GM + transformation lead	One 7-day ship list; public owner commitments
Weekly	Commercial Council: translate decisions into route-to-market execution, one play at a time	30 min	Commercial / key-account owner	The one play of the week, with participation tracked
Weekly	Growth Review: experiment results vs. pre-registered criteria	45 min	Growth owner	Next sprint's test slate; kill / scale / iterate calls; velocity metrics updated
Per request (SLA-bound)	Spend & Promotion Gate: margin integrity, forward-buying risk, compliance	As needed	Finance + compliance	Approve or reject with a 60–90 day true-ROI projection
Monthly	Monthly Business Review: trajectory vs. driver tree, blockers, reallocation, runway and concentration covenants	3 hours	Executive team	Reallocation decisions; kill list; covenant attestation
Quarterly	Strategic Reset: offer architecture, spend strategy, relationship development, pipeline convexity audit	Half day	CEO	Updated capital allocation; threshold-progress review
Annual	Strategic Offsite: portfolio convexity audit, network topology, threshold planning, survivorship stress test (40% revenue-decline scenario)	2 days	CEO / Chairman	Annual capital allocation framework; threshold targets

Two non-obvious properties. First, every forum exists to produce exactly one decision artifact — a ship list, a play, a kill call, an allocation — and a forum that produces only shared awareness is deleted from the calendar. Second, the stack is bidirectional: information flows up as scorecard movement, escalations and tripwire flags; decisions

flow down as ship lists, the weekly play, reallocations and annual capital allocation. A cadence that only moves information up is surveillance. A cadence that only pushes decisions down is proclamation.

OPERATING DOCTRINE

Read the stack as a control system, not a reporting hierarchy. The daily layer is the sensor array. The weekly layer is the actuator. Monthly and quarterly are the feedback controllers that retune the system. The annual offsite is the only forum permitted to change the system's objectives. When an organisation complains it has too many meetings, the diagnosis is almost never the count — it is that decisions are being made at the wrong layer, so every forum re-litigates the layer below it. A ship list is never debated at the MBR; capital allocation is never debated at the huddle.

C.2 Meeting Physics

Four rules keep the stack lean. They are rules of physics, not etiquette.

Rule	Standard	Enforcement test
	The war room reviews the turnaround scorecard. The MBR reviews the driver tree. The Growth Review reviews the experiment slate. Never merged.	A meeting with two scorecards is two meetings badly merged, and the merged version always defaults to whichever scorecard the senior person prefers.
	Status lives in dashboards that systems and agents maintain. Humans convene only to decide. Any agenda item that cannot name its decision output in advance is removed and routed to a dashboard.	If the meeting ends without a change to a ship list, an allocation or a kill list, it was an update meeting and does not recur.
	Every lever, experiment and initiative has exactly one accountable name.	When two executives co-own a relationship, it gets attention exactly when both are free — which is never at the moment the decision is being made on the other side.
	New initiatives enter only through a standard intake form, and only by displacing something on the stop-doing list. Default limit: five.	When nothing can get approved, everything gets proposed, and the backlog becomes a shadow portfolio consuming management attention without producing a single shipped outcome.

C.3 Forum Run-Sheets

Each run-sheet is the complete instruction for the person chairing. Chair changes should not change the meeting.

Morning Protocol — 45 minutes, individual, before the first reactive commitment

Block	Time	Instruction
Survivorship check	10 min	Cash position, yesterday's demand or order data, any tripped tripwires, any covenant proximity. The question is never "how are we doing" but "did anything move yesterday that changes what we must do today."
Convexity prioritisation	20 min	From everything competing for the day, select the single action with the most asymmetric payoff and time-block it before the calendar fills with symmetric, low-upside obligations. One action, not three. An executive completing one high-convexity action per working day compounds roughly 250 per year; the executive dispersed across fifteen reactive threads completes near zero.

Block	Time	Instruction
Bridge-building contact	15 min	One outbound contact to a cluster where the business has a structural hole: a capital source, a buyer at an account where you are thin, a regulatory or trade body, a technical community adjacent to your claims. One contact a day is 250 network edges a year built deliberately rather than accidentally.

Executive War Room — 30–45 minutes, weekly

Element	Standard
Inputs	One scorecard. Last week's ship list with binary status. Tripwire register.
Rule of the room	Review deltas, not levels. A metric that did not move receives no airtime. A metric that moved without an owner who can explain why gets a named investigation due before the next session.
Standing question	When any two tripwires fire in the same quarter: are we buying today's volume with tomorrow's baseline erosion? The burden of proof shifts to whoever wants to continue the current posture.
Output	One 7-day ship list with public owner commitments. Binary — shipped or not.
Enforcement	Public commitment. Executives protect their hit rate in front of peers more reliably than any reporting structure compels them to.
Never discussed here	Capital allocation. Strategy. Anything the MBR or Reset owns.

Commercial Council — 30 minutes, weekly

Element	Standard
Purpose	Translate war-room decisions into route-to-market execution, one play at a time.
Failure mode it exists to prevent	The simultaneous rollout — the quarter where the front line is asked to push a new hero offer, execute an overlay, open a new channel and fix a service problem at once, and does all four at 60% quality.
Output	One play. Distribution push in a named account, an execution standard, an upsell script — with participation tracked by account and by rep.
Operating metric	Participation rate on the current play. A play with adoption below 80% after two weeks is redesigned, not re-announced.

Growth Review — 45 minutes, weekly

Element	Standard
Inputs	Every active experiment with its signed pre-registration doc and its measured result.
Verdicts	Three only: kill, iterate, scale. "Keep watching" does not exist — it is how zombies consume budget. A test that reached its required sample and missed its gate is killed in the meeting.
Velocity accounting	Update four numbers every week: tests shipped, cycle time per test, cost per test, win rate. Most companies cannot answer any of the four.
Health check	A pipeline scaling everything it tests is not discovering wins; it is funding narrative. Expect roughly one in three well-run tests to improve its target metric, and hold a kill rate at or above 50%.

Element	Standard
Output	Next sprint's test slate, with pre-registration docs signed before the meeting ends.

Spend & Promotion Gate — as needed, SLA-bound

Element	Standard
Who sits	Finance and compliance together. The requesting business owner presents; they do not vote.
Three tests	Margin integrity — does net contribution stay positive over the full cycle? Forward-buying risk — is this borrowing volume at a discount and calling it growth? Compliance — claims, pricing representations, partner commitments.
Evidence standard	A 60–90 day true-ROI projection measured against baseline, not against the promoted peak. Spike minus discounts, returns, deductions and post-event decay.
Rejection standard	Every rejection states the condition under which the proposal returns. A rejection without a return condition is a relationship problem waiting to happen.
Why this is the highest-ROI meeting in the stack	Discretionary trade and promotional spend is routinely the second-largest line on the P&L at 15–25% of gross revenue, roughly 59% of promotions lose money globally and 72% in the US, and only about 22% of companies can measure the spend at event level at all.

Monthly Business Review — 3 hours

Element	Standard
Scorecard	The driver tree. Trajectory against each Level-1 and Level-2 node, decomposed by cohort.
Standing agenda	Covenant attestation by name (runway, concentration, position sizing). Reallocation decisions. Kill list. Pipeline convexity scoring. Revenue per FTE and automation coverage.
Reinvestment test	Every saving identified in the month names its convex destination in the same meeting. No destination, no credit for the saving.
Output	Reallocation decisions and a signed covenant attestation — not a review of what happened.

Quarterly Strategic Reset — half day | Annual Offsite — two days

Forum	Mandatory agenda items	Output
Quarterly Reset	Convexity audit of the entire initiative pipeline — eliminate anything without genuine multiple-return upside. Threshold-progress review. Hero-offer and pricing architecture. Relationship development plan. Any proposed change to the operating system itself.	Updated capital allocation and a revised WIP slate. This is the only mid-year forum permitted to change a rule in this manual.
Annual Offsite	Portfolio convexity audit. Network topology mapping and structural-hole plan. Threshold planning. Survivorship stress test against a 40% revenue-decline scenario. Objectives for the operating system itself.	Annual capital allocation framework and threshold targets. The only forum permitted to change the system's objectives.

C.4 Decision Rights — Week-One Infrastructure

Speed is impossible without pre-cleared authority. Publish this grid in week one, adapted to your nouns. Decision rights that live in the CEO's head still route every decision through the CEO's calendar.

Decision right	Responsible (one name)	Consulted (capped)	Informed	Latency standard
Offer / portfolio authority — park items, define hero architecture, price-pack structure	Chief commercial / brand owner	Finance, supply chain (≤ 2)	Executive team	≤ 5 days
Promotion authority — depth, frequency, mechanics	Revenue / sales owner	Finance via the Spend Gate (≤ 2)	Key-account leads	≤ 3 days
Incentive / overlay authority — launch, cap, sunset	CFO delegate + Gate	Compliance (≤ 1)	War Room	≤ 3 days
Claims and compliance language	General counsel / regulatory head	Brand owner (≤ 1)	Marketing, enablement	≤ 5 days
Customer and partner communications	CMO / communications owner	Category owner (≤ 1)	Executive team	Same day for crisis; ≤ 2 days otherwise
Headcount requisition	CEO	Function head (≤ 1) after a failed automation review	Executive team	≤ 10 days; no path without the automation review on file
Vendor commitment above materiality threshold	Function head	Finance, security (≤ 2)	Executive team	≤ 10 days; no signature without matrix, pilot result and pre-mortem

Two features matter more than the assignments. The latency standard converts "finance is slow" from a grievance into a measurable miss — a decision right without a service-level commitment is a suggestion. And the cap on Consulted makes the cost of each added approver explicit: left unmanaged, consultation lists grow one name per organisational embarrassment, each addition rational in isolation and collectively ruinous.

FROM THE FIELD

A \$120M business carried 38 active "strategic initiatives" and an 11-week average approval latency for any spend decision above \$25K, on a CEO calendar holding 34 hours a week of recurring meetings with no defined decision outputs. Installing the stack took one week: WIP limit of five, a standard intake form requiring every new initiative to displace one on the stop-doing list, a weekly war room with a single scorecard, and a Monday ship list with named owners. Within eight weeks approval latency fell to 9 days, initiative count fell to seven, and the spend gate killed two auto-renewed negative-ROI programs — recovering roughly 180 basis points of gross margin before a single new growth initiative shipped. Nothing about the strategy changed. The physics changed.

PART D

The P&L Spine

One accounting of reality. Every metric on every scorecard must trace through an explicit driver tree to a P&L or cash-flow line. A metric that cannot be traced is deleted; a meeting that reviews untraceable metrics is cancelled. This is the discipline that separates data-driven operations from dashboard theatre.

The alternative is present in every stalled business: forty metrics across eleven dashboards, none reconciled, each owned by a function that picked the definition flattering its own quarter. Sales reports gross revenue; finance reports net revenue a third lower; marketing reports lift on a seven-day window while the baseline erodes in week two. Nobody is lying. Everybody is wrong.

D.1 Construction Rules

The tree is business-model agnostic. Revenue is always customers $\times$ behaviour $\times$ value per behaviour. Margin is always price realised minus the cost of delivery and the cost of the incentives used to move it. Three properties make a tree operational rather than decorative.

Property	Requirement	Why it is enforced
	Level-1 drivers multiply back to the P&L line with no residual "other" bucket.	A residual is where unexamined economics hide.
	Every Level-2 metric has exactly one accountable name and one owning forum reviewing it on a fixed cadence.	An unowned metric drifts and nobody is accountable for saying the sentence out loud.
	Every Level-2 metric is computable per cohort.	The average is where dying economics hide. Any metric that materially changes when decomposed is managed only in decomposed form.

OPERATING DOCTRINE

Averages are a reporting convenience, not an operating reality. Any executive presenting an average without its cohort decomposition has not finished the analysis. Cohort discipline is the difference between a correctable drift and a public restatement.

D.2 The Canonical Tree

P&L line	Level-1 drivers	Level-2 operating metrics (scorecard level)	Owning forum
Revenue	Active customers $\times$ purchase frequency $\times$ AOV; in retail: weighted distribution $\times$ velocity $\times$ price/mix	Weighted (%ACV) distribution; total distribution points; velocity per unit of distribution; first-purchase activation rate; 30-day repeat-purchase rate; reactivation rate; subscription penetration; key-account participation %	Weekly War Room
Gross margin	Price/mix, discount depth, returns and deductions, freight	True promotion ROI (60–90 day net contribution); gross-to-net waterfall position; returns % by item and by event; discount-depth trend; freight per unit	Spend & Promotion Gate

P&L line	Level-1 drivers	Level-2 operating metrics (scorecard level)	Owning forum
Sales & marketing	CAC by channel; discretionary spend efficiency	CAC by channel; trade/promotional spend as % of gross revenue (15–25% benchmark band); inbound share of pipeline; LTV:CAC by cohort (3:1 floor)	Monthly Business Review
Opex / SG&A	Headcount, software, agencies, logistics	Revenue per FTE; cost per experiment; agent-automation coverage of recurring workflows; vendor TCO vs. matrix score	MBR / Quarterly Reset
Cash	Cash conversion cycle; capex; deployments	13-week rolling cash forecast; months of runway; concentration (largest exposure / net revenue); Kelly-sized position audit	Monthly covenant review
Enterprise value	Convex position portfolio; reputation assets; threshold progress	Convexity-weighted pipeline value; documented wins toward the 2–3 threshold; reference-network growth	Quarterly / Annual

Three design decisions deserve comment. Revenue metrics are reviewed weekly because revenue drivers move weekly and are correctable weekly; cash and enterprise value move slowly and are reviewed monthly and quarterly, which prevents re-litigating strategy in a Monday meeting. Discretionary spend sits under S&M as a percentage of gross revenue but its ROI is gated under gross margin — the spend is a commercial decision, the truth about the spend is a margin fact, and separating the two is what breaks the sales-versus-finance incentive mismatch that lets unprofitable programs repeat annually. And the enterprise-value row is not decorative: a business that manages only this year's P&L reliably consumes its own future.

D.3 Instantiations — Same Spine, Different Clothes

Business model	Level-1 decomposition	The metric that decides everything	The concentration tripwire
Retail / CPG brand	Weighted distribution × velocity per door × price/mix	Velocity per point of distribution. Distribution can be purchased; velocity must be earned.	Largest retailer as % of net sales
DTC / subscription	Active customers × purchase frequency × AOV	30-day repeat-purchase rate by acquisition cohort	Largest acquisition channel as % of new customers
Multi-unit / hospitality	Units open × average unit volume × same-unit growth, decomposing to traffic × average check × frequency	Unit-level contribution margin by vintage and format	Single geography or format as % of contribution
Commercial lending / credit	Loan book × net interest margin × fee attach rate	Credit losses measured against the vintage of the cohort that generated them, not this quarter's originations	Single-counterparty exposure vs. capital
Professional / advisory services	Active engagements × realisation rate × average engagement value	Utilisation of senior judgment hours vs. delegated hours	Largest client as % of revenue

CONSEQUENCE FRAMING

When a financial intermediary cannot reconcile its own liabilities to a single source of truth, the failure is usually reported as a technology failure. It is not. It is a driver-tree failure: no single decomposition existed that every forum could trust. The 2024 collapse of one banking-as-a-service intermediary froze more than \$150 million

belonging to roughly ten million consumers across dozens of downstream apps — this chapter's opening warning, realised at consumer-harm scale.

D.4 The Tripwire Register

Tripwires are not trends. Each is a pre-set threshold whose breach changes the default action without further debate. Publish this register, assign every row an owner, and automate the detection.

Tripwire	Threshold	What it means	Default action on breach
Creeping discount depth	Depth trending up two consecutive quarters	Margin is being refinanced as volume. Customers are being trained to wait.	Auto-hold on any new event above current average depth until the Gate re-underwrites the calendar.
Rising returns and deductions	Above the 5–15% of gross benchmark band, or any single item above 15%	A channel-economics problem masquerading as an accounting one.	Item-level review at the next MBR; suspend expansion in the affected channel.
Accelerating early-cohort cancels	Months 1–3 cancel rate rising vs. prior cohorts	The acquisition recruited deal-seekers, not customers. The DTC equivalent of forward-buying.	Freeze the acquisition source that produced the cohort; re-read the onboarding experiment.
Declining cohort LTV	Any decline by vintage on a gross-margin basis; LTV:CAC below 3:1	Acquisition quality, not retention effort, is the problem. LTV is hypersensitive to churn — a drift from 5% to 8% monthly cuts subscriber lifetime value by roughly 40%.	Reallocate acquisition spend by cohort quality, not by volume, before the next flight.
Spend escalation against a falling baseline	Discretionary spend rising while non-promoted velocity falls	The signature of promotion addiction: each event must be deeper because the customer has been trained to wait.	Mandatory standing question in the War Room; burden of proof shifts to whoever wants to continue.
Concentration drift	Any single customer, channel or vendor above 20% (flag) or 25% (breach)	A single planogram reset, contract loss or API deprecation becomes a company-level event.	Named diversification plan with dates; no new concentration-increasing commitments.
Runway compression	Below 12 months on the 13-week rolling forecast	The clock that funds patience is running out. Forced exits from positions that needed only time.	Auto-hold on outside-zone deployment. Reset the bet portfolio first.
Quality or compliance incident	Any open severity-1	Reputation is a capitalised asset. This is an impairment event.	Stop-the-line in the affected workflow. Nothing ships until closed.

STOP CONDITION

When any two tripwires fire in the same quarter, the standing question becomes mandatory in the War Room: are we buying today's volume with tomorrow's baseline erosion? The burden of proof shifts to whoever wants to continue the current posture, and the evidence standard is the 60–90 day true-ROI window — spike minus discounts, returns, deductions and pull-forward decay — not the lift chart.

D.5 The Hidden-Economics Briefing

Three facts every operator in the business is expected to know without looking them up. They define the terrain the tree is built on.

Fact	The number	The operating consequence
	Trade and promotional spend 15–25% of gross; partner deductions 5–15%; returns and allowances 3–10%. Net revenue lands at 60–70% of gross.	Anyone quoting gross is overstating the top line by roughly a third, and every downstream ratio computed on gross is flattered by the same amount.
	Roughly 59% of promotions lose money globally and 72% in the US. Of every \$100 of promotional sales, ~\$35 subsidises buyers who would have paid full price and ~\$8 cannibalises your own full-price volume. Post-event velocity runs at 40–60% of baseline in week one, with baseline restored only around week four — six weeks at deep discount.	Any event judged on lift-week sales is measuring a loan, not a gain. One brand adding decay to its model watched measured ROI fall from 1.8× to 0.7× — it had been losing money while believing it was profitable.
	Worldwide out-of-stock has sat at roughly 8% — one item in thirteen — since it was first measured in 1996 and reconfirmed across 32 countries. It costs retailers roughly 4% of sales, and about 30% of out-of-stock items are bought from another store.	Thirty years of technology investment did not move the base rate. Execution systems do. Carry invariants as fixed costs of reality, not as improvable metrics.

D.6 The Cohort Standard

Aggregates are the enemy of truth. Every hidden-economics failure above is invisible in the average and obvious in the cohort. The minimum segmentation standard:

- Customers by lifecycle stage (0–30, 31–90, 90+ days) and by lapse status (30–90, 90–180, 180+ days).
- Channel or account performance by velocity cohort, so that a headline distribution number can never again conceal a velocity collapse in the largest account.
- Subscribers by acquisition-month cohort, with cancel curves tracked separately for the first 90 days.
- Programs and promotions judged as cohorts, with each event's decay curve compared against its own category base rate.

Why the 30-day repeat rate is the highest-information metric you own: across a panel of 156,110 DTC customers, 18.8% placed a second order within 365 days — but of those who repeated, 50.3% did so within 30 days and 76.4% within 90. A customer who makes a second purchase is roughly 45% more likely to make a third, and a 10-point repeat-rate lift typically drives a 25–40% increase in lifetime value. Repeat behaviour is decided in the first month, and it is only visible per cohort.

Reactivation is the cheapest revenue in the system: win-back on recency-segmented lapsed cohorts typically recovers 8–15% of the lapsed base at roughly one-third of blended acquisition cost, and reactivated customers carry 1.3–1.8× the LTV of new acquisitions — but only if lapse cohorts are tracked finely enough to intervene before the 180-day cliff. A team buying new acquisition while ignoring the lapsed file is paying triple for worse economics.

FROM THE FIELD

A business celebrating 92% weighted distribution and record quarterly shipments had, on one afternoon of tree-

building, lost velocity from 6.1 to 4.2 units per store per week over eight quarters, seen trade spend climb from 18% to 27% of gross, and was sourcing 61% of promotional volume from two deep-discount events whose post-event weeks ran below baseline. The record quarter had been purchased. Distribution expansion was frozen, the two loss-making events were cut, and the freed spend went behind velocity programs in the 40% of doors still performing. Net revenue fell 4% for two quarters — then grew 11% the following year on 2,000 fewer doors.

PART E

The Growth Engine

The experiment pipeline is the production line of the entire system. The driver tree, the cadence stack and the gates exist to feed it, fund it or enforce its verdicts.

The statistical fact that justifies all of it: roughly one third of ideas tested at Microsoft-scale experimentation programs improved the metric they were designed to improve; about a third did nothing and a third actively harmed it. At Google, roughly 10% of some 12,000 experiments in a single year led to actual business changes. Netflix operates on a standing assumption that 90% of what it tries is wrong. Read those together: expert intuition — yours, mine, your CMO's, your board's — is a coin flip weighted against you. Testing velocity is not a growth tactic. It is the only honest epistemology available to an operator.

CONSEQUENCE FRAMING

A business shipping four major initiatives a year on conviction alone will, at the observed base rate, spend roughly two-thirds of its growth budget on initiatives that do nothing or do harm — and will never know which third worked, because nothing was measured against a control. Across a five-year plan, that is the difference between compounding and drift.

E.1 Pipeline Architecture

Treat it as a factory: raw ideas in, scaled standard work out, with scrap rates measured at every station.

Station	Mechanic	AI role	Failure it prevents
1. Backlog	Every idea enters one backlog — no side channels, no executive favourites jumping the queue — scored on convexity-adjusted ICE.	Draft scoring, dedupe, surface cohort anomalies that generate new candidates.	The shadow portfolio: proposals that consume management attention without ever shipping.
2. Pre-registration	A one-page experiment doc written and signed before launch.	Power calculations, sample-size math, guardrail selection.	Post-hoc rationalisation. Un-pre-registered tests produce the same failure rate but convert failures into anecdotes that get scaled anyway.
3. Execution	Bounded audience, correct randomisation unit, monitored for sample-ratio mismatch.	Variant generation, QA checklists, mismatch monitoring, guardrail alerts.	Contaminated tests that poison downstream decisions.
4. Weekly verdict	Kill, iterate or scale against pre-registered criteria. No fourth verdict.	First-pass readouts within 24 hours of significance.	Zombies. "Keep watching" is how a failing test consumes budget indefinitely.
5. Standard work	Scaled winners enter Phase 5 of the loop: codified, trained, and into the play of the week.	Playbook drafting, training asset generation.	The next regime change quietly reversing a win nobody wrote down.

The kill branch carries roughly two-thirds deliberately. A pipeline whose kill rate runs below half is either not testing real ideas or not enforcing its gates — both are defects. The failure archive is an asset: each killed test is purchased information about the driver tree that no competitor can see.

E.2 Convexity-Adjusted ICE Scoring

Component	Scored on	Anchor for a high score
Impact	Projected P&L movement traced through the driver tree in one step.	Moves a named Level-2 metric with a stated mechanism, not a narrative benefit.
Confidence	Base-rate evidence, not enthusiasm.	A defensible outside view exists: what happened the last ten times a company like ours tried this.
Ease	Cost and cycle time to a readable verdict.	Readable inside 14 days at a budgeted cost.
	Whether the upside is structurally unbounded.	Pricing, offer architecture, channel unlocks and subscription design outrank bounded polish. A pricing win compounds on every future order; a copy win is a one-time step change.

Portfolio standard: more than 60% of the active slate must be structurally unbounded. Deliberately barbell the slate — high-convexity plays whose wins compound on every subsequent cohort, plus a smaller set of structured bets with explicit kill economics. What is absent is as important as what is present: no packaging refreshes, no awareness campaigns, no brand-health initiatives. Every test traces to a driver-tree node in one step.

E.3 The Pre-Registration Form

One page. Signed before launch. Read aloud at the gate. If a field is blank, the test does not launch.

Field	Standard
Hypothesis	Falsifiable form. "If we do X, then [named metric] will move by at least Y, because [mechanism]."
Target metric	Exactly one, named on the driver tree. Metric shopping is a named defect.
Minimum detectable effect	Stated as a number before launch. If the MDE is smaller than the effect you would act on, the test is not worth running.
Sample size and run time	Computed from the MDE. No peeking-driven early stops.
Randomisation unit	Customer, store cluster, market or session — stated explicitly, with a sample-ratio-mismatch check scheduled.
Guardrails	Trust, margin and returns at minimum. A win on the target metric that breaches a guardrail is a loss.
Holdout	For any always-on program, a persistent global holdout. Without it, reported ROI is partially fiction.
Kill criteria	The number below which the test dies on its scheduled date. Written before the data arrive.
Decision map	What we do if it wins, what we do if it loses, what we do if it is inconclusive. All three, in advance.
Signatures	Test owner and the forum chair. Two names.

E.4 The Six Statistical Non-Negotiables

Experimentation without statistical discipline is astrology with dashboards. Violating any one converts the pipeline back into a conviction machine while retaining the costume of rigour.

Discipline	Rule	Failure mode prevented
Sample size and power	Compute MDE and required N before launch; no peeking-driven early stops.	Declaring victory on noise — the false-positive factory.
Cohort integrity	Randomise at the correct unit; check sample-ratio mismatch.	Contaminated tests that poison downstream decisions.
True-ROI window	Promotions and spend events judged on 60–90 day net contribution including post-event decay.	Promotion addiction; buying volume with tomorrow's sales.
Base-rate check	Every forecast compared to the outside view before approval.	Planning-fallacy budgets; hasty generalisation from pilot euphoria.
Holdouts	Maintain global holdouts for always-on programs.	Claiming credit for revenue that would have arrived anyway.
One-metric focus	Each test moves one pre-registered driver-tree metric; guardrails on trust, margin and returns.	Metric shopping; silent harm to brand and contribution margin.

E.5 Velocity Accounting

The master metrics of the pipeline itself. Most companies cannot answer any of the four. Update them weekly.

Metric	Target	Interpretation
Tests shipped per week	≥ 3	The only lever on expected wins when the base rate is fixed is the number of shots.
Cycle time per test (idea → verdict)	< 14 days	Cycle time is the cadence stack showing up in the pipeline.
Cost per test	Falling 2–5× against baseline	When cost per test falls 5×, the rational number of concurrent tests rises 5× at flat opex — and base-rate math converts that directly into wins per quarter.
Win rate vs. base rate	$\approx 33\%$, kill rate $\geq 50\%$	A win rate far above base rate means the tests are trivial or the gates are not being enforced. Both flatter the number.
Convexity ratio of the slate	> 60% structurally unbounded	Below this, the pipeline is polishing rather than positioning.

OPERATING DOCTRINE

Run the experimentation system like a production line: one backlog scored on convexity-adjusted ICE; every test pre-registered with MDE, sample size and kill criteria; velocity accounting on tests shipped, cycle time, cost and win rate; three verdicts only — kill, iterate, scale; winners codified into standard work. Expect two-thirds of everything to fail, and build the economics so the surviving third pays for all of it.

E.6 Experimentation Without a Browser

The honest objection: a digital marketplace can run a thousand concurrent tests because it owns a website; a physical business cannot A/B test a shelf. The objection is half right and wholly defeatist. What a physical business has instead is an experimentation continuum.

Mechanism	What it is	The gate it replaces
Telemetry as live test	Instrumented points of consumption generating millions of micro-decisions per day. One beverage platform runs 50,000+ connected dispensers pouring roughly 11 million drinks daily; products have reached national retail because the mixing data said so, cutting data-to-shelf time to as short as 90 days against a traditional 18-month innovation cycle.	Concept testing and focus groups.
Channel sequencing as staged experimentation	Prove per-unit productivity in one channel before buying the next. One challenger reached roughly \$200M in revenue on ~28,000 doors where a typical brand needs 80,000, by sequencing natural → grocery → mass and proving velocity per door at each stage.	National launch on conviction.
Staged geography	Test markets with matched controls and pre-registered per-unit floors, expanded ring by ring.	Simultaneous rollout.
Pilot-first procurement	A paid, bounded trial with pre-registered adoption metrics before any annual contract.	The demo-plus-quarter-end-discount purchase.

The arbitrage is the point. When an industry has no published experimentation base rate, the first operator to install a real pipeline is not playing catch-up — they are playing alone.

PART F

Talent

A senior hire is the most asymmetric position most operating portfolios ever hold. The downside — salary, severance, lost quarters, team drag, replacement cycle — is conservatively 5–15× annual compensation for senior roles. The upside compounds for years across every metric the person touches. Under a lean-headcount doctrine, every hire is rare, senior and enormously leveraged: when one person with agents and workflows does the work of a five-person team, each seat is a concentrated position that must clear the same evidentiary bar as a capital deployment.

The base rates justify the rigour. A century of selection research is consistent on the rank ordering across every re-estimation: general mental ability, work samples and structured interviews beat conversation, credentials and tenure, every time. Unstructured interviews — the default method of most companies — sit near the bottom at an operational validity of roughly .38, while structured interviews and ability measures sit at the top and combinations dominate everything else. Years of experience (.16–.18) and years of education (.10) barely move the needle. Hiring on charisma and conversation is the selection-science equivalent of investing on narrative.

F.1 The Hiring Loop

Phase	Action	Standard
	Write the outcome scorecard, not the job description.	Three to five measurable outcomes this role must deliver in 12 months, each traced to a named node on the driver tree, plus the behavioural competencies that predict them. Job descriptions describe activities; scorecards describe results.
	Structured process, pre-committed rubric.	Same questions, same order, same anchored 1–5 rubric, every candidate. Anchored means each score point is defined by observable behaviour, not adjectives. Score independently, in writing, before any group discussion.
	The paid work-sample test.	Every finalist performs a paid, time-boxed sample of the actual job on real or realistically anonymised material. Pay market rates; paid work is real work and keeps the exercise legally clean. Blind-score where feasible.
	Evidence-weighted references.	Arrange references through the candidate: "please set up a call with your manager from that role." The arranging behaviour is itself a signal. Ask outcome-anchored questions and triangulate — a claim surviving three independent references earns its weight; one surviving an enthusiastic champion and two awkward silences does not.
	The 90-day gate.	Every offer includes a written 90-day plan with pre-registered success criteria agreed before the start date. Day 90 is a real gate with three verdicts: scale, iterate (one 60-day recheck only), or kill.

OPERATING DOCTRINE

Every hire is a pilot with a day-90 gate. The offer letter is not the decision; it is the start of the final validation leg. A hire you would not make today, with full information, has already failed — the only question is whether you pay the exit cost now or the compounding cost for the next three years.

F.2 The Outcome Scorecard Template

One page, written before sourcing begins. Instantiate the nouns for your role; the structure does not change.

Outcome (driver-tree node)	Baseline	12-month target	Measured by
[Level-2 revenue metric, e.g. velocity per unit of distribution]	[current]	[target]	Panel or POS data, monthly
[Activation or conversion metric]	[current]	[target]	Funnel analytics, weekly
[Retention metric, e.g. 30-day repeat rate]	[current]	[target]	Cohort dashboard
[Margin metric, e.g. true spend ROI]	[current]	[target]	Post-event decomposition
[Pipeline or coverage metric]	[current]	[target]	CRM, stage-weighted

Competency block (four standing entries): quantifies instinctively; designs experiments with pre-registered metrics; runs the weekly review to standard work; multiplies own output with tooling and agents.

F.3 The Candidate Evaluation Matrix

The matrix makes the decision mechanical precisely where persuasion pressure is highest. Senior candidates are, by definition, professionally persuasive. Weights are pre-committed, evidence sources fixed, anchors written before the first interview.

Dimension	Weight	Evidence source	Anchor for a 5
Outcome track record vs. scorecard	30%	Structured deep-dive on last 2–3 roles; quantified results, reference-triangulated	Delivered analogous outcomes with documented numbers in comparable constraint environments
Work-sample performance	25%	Paid, time-boxed real-work exercise, blind-scored where feasible	Output usable tomorrow with minor edits; reasoning transparent, data-anchored, assumptions stated
Problem-solving / cognitive signal	15%	Case reasoning within the structured interview	Decomposes novel problems to first principles; quantifies instinctively; changes position on evidence
AI-native leverage	15%	Demonstrated multiplication of own output through tools, agents and designed workflows	Operates at demonstrably > 2× conventional throughput for the role; builds workflows others reuse
Values & decision hygiene	10%	Behavioural questions on ethics, dissent and loss-admission; reference triangulation	Has publicly reversed course on evidence; names own past errors with root causes, not narratives
Network / structural-hole value	5%	Cluster map of the candidate's reachable networks vs. the firm's current gaps	Bridges a cluster the firm does not currently reach — capital sources, buying organisations, regulatory bodies, technical communities

The two mechanical rules

- A weighted composite below 4.0 is a no-hire, regardless of enthusiasm in the room. The empty seat is cheaper than the wrong occupant: a mis-hire's true cost is conservatively 5–15× annual compensation, while the vacancy cost is measurable and bounded.
- No interviewer sees another's scores before committing their own. Independence is what makes the panel five measurements instead of one. Debrief-first processes reliably convert a panel of five interviewers into one interviewer with four witnesses.

CONSEQUENCE FRAMING

Skip the matrix once and you have set the precedent that the matrix is optional. The first charismatic exception that gets hired teaches the organisation that process yields to persuasion — and persuasion is precisely the variable the process exists to control. One mis-hired executive costs the 5–15× directly; the precedent costs more, because the next three hiring loops run on narrative again and the losses compound off the books.

F.4 The Performance Cadence

Annual reviews are batch processing applied to the one asset that most needs continuous flow. Latency in feedback is latency in correction, and latency in correction is money.

Cadence	Instrument	Standard
Weekly	Public commitments in the War Room	Each owner states next week's ship list and reports on last week's. The commitment is public, the miss is public, and the pattern is visible to everyone within a month.
Monthly	One-on-one against the role scorecard	30–60 minutes with the scorecard on the table. Three questions: what moved, what is stuck, what decision do you need from me. The conversation is against the role, not against the manager's mood.
Quarterly	Written scored evaluation	Every manager writes a scored evaluation of every report on the same anchored rubric used at hiring. Identical instrument end to end means the performance file becomes longitudinal data rather than an annual essay contest — and you can ask, with numbers, whether your hiring signals predict performance.

The four-dimension performance rubric

Dimension	Weight	Anchor 5	Anchor 1 (auto-escalation)
Outcomes vs. role scorecard	40%	All Level-2 metrics at or above plan with evidence of mechanism, not luck	Metrics missed two or more quarters with narrative explanations instead of mechanisms
Decision quality	25%	High gate hit rate; pre-mortems filed and predictive; checklists followed; losses admitted fast with root causes	Repeat gate failures; decisions reconstructed after the fact; no decision-journal entries
Leverage / multiplication	20%	Output per unit of resource demonstrably > 2× role norm; builds workflows and agents others reuse	Throughput below norm; treats tools as optional; output does not scale with resources

Dimension	Weight	Anchor 5	Anchor 1 (auto-escalation)
Trust behaviour	15%	Communication hygiene impeccable; flags bad news early; compliance record clean; dissents openly then commits	Late bad news, filtered reporting, or any compliance event

Why outcomes carry 40% and not 100%: in a promotion-heavy quarter a manager can hit the number by buying volume with unprofitable spend. Decision quality and leverage catch exactly that failure mode by asking how the number was made. Decision quality is scored on artifacts — gate memos, pre-mortems, journal entries — not on the manager's memory of how smart the person sounded in meetings. Artifacts do not have charisma bias.

F.5 The Keeper Test and Separation Protocol

The most expensive inventory in a business is not aging stock in the warehouse; it is the unresolved people decision sitting in a manager's head. The keeper test converts a vague, socially expensive feeling into a discrete, answerable question with a timestamp.

Step	Rule
The question	"If this person told me they were leaving for a comparable role elsewhere, would I fight hard to keep them?" Answered quarterly, in writing, per report, with two sentences of justification tied to the rubric scores.
A "no" triggers diagnosis, not termination	One "no" quarter starts a structured intervention: a specific gap named against the scorecard, a written improvement plan with 90-day checkpoints, and honest conversation. Many first "no"s reverse — the feedback was simply never delivered clearly before.
Two consecutive "no"s is a decision already made	At that point the organisation is not deliberating; it is refusing to execute. The two-cycle rule exists because humans are excellent at converting "not yet" into "never" when the decision is uncomfortable.
Execute within two weeks	Delay after the decision is pure cost. The departing employee senses it, the team senses it, and the manager burns credibility explaining a decision already made.
Severance above market norm	Generosity is cheap insurance on two fronts: every separated employee becomes a permanent narrator of how your company treats people, and litigation risk — which threatens the survivorship covenants protecting the whole system — drops sharply when the exit is dignified and the package visibly fair.
Honest without punitive	The message is fit and trajectory, delivered with respect. Public humiliation serves no operating purpose and poisons the narrator effect.
Post-mortem filed	Every separation is also a hiring-process defect. Which interview signal failed, which reference was skipped, which scorecard dimension was waved through. Feed it back into the rubric.

CONSEQUENCE FRAMING

Refusing to execute carries three compounding costs. The seat's metrics run below plan for every additional quarter. Top performers calibrate their own effort to the standard the organisation visibly tolerates — nothing erodes an A-player's engagement faster than carrying a gap while leadership pretends not to notice. And every month of inaction broadcasts that the stated performance bar is fiction; organisations believe what they observe, not what the values poster says. A "no" held for two years is not a personnel problem. It is a three-front cultural compounding loss.

F.6 Development as Capital Allocation

Development spend follows power-law logic, not egalitarian logic. No executive would split an innovation budget equally across twenty products regardless of velocity; splitting the development budget equally across twenty managers regardless of compounding rate is the same error in a less visible ledger.

Tier	Who	Allocation
	Everyone. Non-negotiable and universal.	Standard work, a written role scorecard, honest quarterly feedback, access to baseline training. Clarity is a right, not a reward.
	Solid performers.	Targeted skill-building against the specific gap the rubric surfaces, on a normal timeline.
	Demonstrated compounders.	Uncapped. Executive coaching, cross-functional stretch assignments, P&L exposure two levels early, seat time in the senior forums, sponsorship for the next role before the next role officially exists.

The discipline cuts both ways. Concentrating investment on compounders is only defensible if the keeper test is actually executed. Otherwise the organisation ends up with the worst possible allocation: development spend spread thinly to avoid hard conversations, and separation decisions deferred indefinitely to avoid harder ones. The egalitarian instinct, applied to talent, is not kindness. It is a refusal to allocate — and refusal to allocate is how companies end up with deep benches of adequately trained underperformers and nobody ready to run the next unit.

PART G

Technology, Vendors and AI

Technology procurement is where Lean discipline and convexity logic meet a hostile counterparty: a professional enterprise sales motion engineered around anchoring, artificial urgency, social proof and authority claims. A bad platform contract does not just waste licence fees — it corrupts the data the War Room uses to allocate the largest discretionary pool in the company.

G.1 Two Rules That Precede Any Scoring

Rule	Standard	Why
	No annual contract precedes a 30–60 day paid pilot with pre-registered success metrics on the driver tree.	Paid, because free pilots are sales theatre — the vendor staffs your pilot with its best solutions engineers and you evaluate a service you will never receive again. Pre-registered, because metrics chosen after the results are in are not metrics; they are narratives. Vendors who refuse a paid pilot with named success metrics are self-identifying on downside risk. Believe them.
	Every contract is scored on the cost of leaving: data export rights, integration unwinding, retraining.	Switching cost is negative convexity — capped upside for the buyer, unbounded rent for the seller. A vendor quoting 20% below market with a three-year lock-in and export fees is not cheaper; it is a lease on your own data with an introductory rate.

FROM THE FIELD

A fitness-apparel brand built its early e-commerce operation on a platform selected without reversibility or peak-load stress-testing. On Black Friday — the single highest-revenue day of the retail calendar — the site went down for eight hours, an estimated £100,000 in lost sales on roughly £41 million of annual revenue, ten months into production. The failure was not a hostage clause or a hidden fee. It was a platform decision made without asking the one question this matrix forces: could this survive its single highest-traffic day?

G.2 The Vendor Evaluation Matrix

Nine criteria, weights fixed before vendor contact, each scored 1–5 against anchored definitions. The weights are doctrine: driver-tree impact carries five times the weight of support quality because a tool that moves a named metric with mediocre support beats a beautifully supported tool that moves nothing.

Criterion	Weight	What a 5 looks like	What a 1 looks like (auto-flag)
Driver-tree impact	25%	Directly moves a named Level-2 metric with a testable mechanism; pilot design agreed in writing.	Benefits described only in narrative — "efficiency," "visibility," "alignment" — with no measurable line to a driver-tree node.
Total cost of ownership (3-yr)	15%	All-in TCO modelled: licence, implementation, integration, training, internal maintenance hours, per-seat escalation at scale.	Quoted price excludes implementation or mandatory modules; opaque consumption pricing that reprices at renewal.

Criterion	Weight	What a 5 looks like	What a 1 looks like (auto-flag)
Integration & API depth	12%	Open, documented APIs; native connectors to the existing stack; agent-accessible endpoints so data flows into automated reporting.	Closed system; manual export only; professional services required for basic connections.
Data ownership & portability	12%	Full export in standard formats at any time, contractually guaranteed; your data trains nothing without written consent.	Data-hostage clauses; export fees; ambiguous training-use language buried in the DPA.
Security & compliance posture	10%	Current third-party attestation; clear data-processing agreement; incident history disclosed voluntarily.	No attestation; evasive on breach history and subprocessors.
Vendor viability & roadmap	8%	Durable financials or strategic backing; roadmap aligned with your three-year architecture.	Runway risk; roadmap driven by the vendor's next funding narrative, not your needs.
AI leverage & trajectory	8%	Genuine model-driven capability that improves with your data volume; agent interoperability.	AI-washing: rules engines and templates rebranded as intelligence; no learning loop.
Reversibility / switching cost	5%	Month-to-month or annual with clean exit; documented migration path; survives its single highest-load day.	Multi-year lock-in, auto-renewal traps, punitive termination; no disclosed peak-load architecture.
Support & partnership model	5%	Named accountability, SLAs with remedies, reference customers at your scale who took real unscripted calls.	Ticket queue only; references curated and rehearsed; SLA without remedy.

Three properties do the real work

- The auto-flag column converts any single catastrophic criterion into a veto. A vendor scoring 1 on data portability is disqualified even with a strong composite, because a hostage clause cannot be averaged away.
- Weights are fixed before vendor contact, so the sales team cannot lobby them upward on the criterion where their product is strongest.
- The "what a 5 looks like" anchors are written against real operating mechanisms, which forces scorers to test claims — "show me the readout on live data," "show me what happens at ten times today's peak load" — rather than accept adjectives. A composite score is a forecast about driver-tree impact, not a mood.

G.3 Scoring Governance

Rule	Standard
Two independent scorers	Both complete the full matrix alone and submit before the evaluation meeting. Divergences of two points or more on any criterion are argued in the meeting, never averaged silently — the disagreement usually locates exactly the ambiguity the vendor engineered.
Thresholds that bind	Composite ≥ 4.0 advances to pilot. Composite 3.0–3.9 advances only with a written mitigation for every criterion below 3. Below 3.0 is a pass, regardless of demo quality, dinners or discount pressure. Any single-criterion 1 is a veto.
Pre-mortem before signature	"It is twelve months from now and this implementation failed. Why?" The three most-cited failure modes — adoption collapse, integration overrun, data-quality mismatch — each get a named owner and a pilot metric designed to surface them early.
Decision journal	Every decision above the materiality threshold enters the journal with its pre-registered pilot metrics, composite scores and pre-mortem. Renewal decisions twelve months later are made against recorded

Rule	Standard
	predictions rather than reconstructed memory.
TCO before quote	Model the full three-year exposure before hearing any number. In a typical evaluation the gap between quoted licence and modelled TCO runs 50%+ — and the size of that gap is itself diagnostic.

Sales-pressure inversions

The enterprise sales playbook runs on three levers. Each inverts into a signal.

Lever	What they are doing	Your inversion
Anchoring	The first number quoted becomes the reference point for every subsequent discussion.	Model TCO before hearing any quote. The first number is a reference, not a price.
Artificial urgency	A discount deadline expiring "this quarter."	Treat it as data about the vendor's quarter, not about your decision. A vendor that discounts 25% to close in week twelve will discount 25% in week one of next quarter.
Social proof	The logo slide.	It tells you the vendor sold to firms like yours, not that firms like yours succeeded. Counter with reference calls you source yourself, at your scale, unscripted.

G.4 Build vs. Buy vs. Orchestrate

The AI era adds a third option to the classic dichotomy — orchestration: thin internal workflow layers over best-of-breed primitives. It also settles the empirical question. Across a review of more than 300 enterprise GenAI deployments, purchased and partnered solutions succeeded roughly 67% of the time versus roughly 22% for internal builds — a 3:1 ratio in favour of buying capability and concentrating internal effort on workflow integration.

Category	Rule	Examples	The waste it avoids
	Commodity capability in mature markets where vendors amortise R&D across hundreds of clients and your team will never out-build them.	Spend-management platforms, syndicated data, demand planning, analytics, core transaction systems.	Building undifferentiated software is overproduction waste — inventory of code nobody outside the building would buy.
	The workflow layers that encode your operating system. This is where bought capability gets converted into P&L impact.	Driver-tree automation joining data sources; the experiment pipeline; War Room scorecard feeds; tripwire monitoring.	The bolt-on failure: tools purchased and never embedded in a real workflow, which is where roughly 95% of enterprise GenAI spend has produced no measurable P&L impact.
	Only where the asset itself is the moat and the training data is unreplicable.	A demand-sensing model trained on your unique panel; proprietary data products; signature IP.	Buying your moat is renting your future.

G.5 The Agent Org Chart

The AI-native enterprise maintains two org charts. The first is short: the accountable humans. The second is long: the agent workflows, each owned by a named human, each carrying a service-level definition, each reviewed in the same cadence stack as any team. The reorganisation the laggards skip is precisely this — drawing the second chart, assigning owners and holding it to service levels.

Agent workflow	Human owner	Service-level definition	Human gate
Scorecard & cohort automation	FP&A lead	Full cohort decomposition refreshed by 6:00 daily; double-firing tripwires flagged.	Every published metric signed.
Experiment operations	Growth owner	Readout within 24h of significance; power calculation pre-registered per test.	Kill / scale / iterate calls only in the Growth Review.
Content & authority production	Growth owner	Five drafts per week per owner, each with a source manifest.	Every external piece, no exceptions.
Pipeline & relationship intelligence	Commercial lead	Key-account briefs Monday 07:00; contact-discipline drafts daily.	All partner-facing sends.
Diligence & document production	FP&A lead	First-draft memo within 48h of request.	All financial and legal outputs.
Customer-facing service & support	Service owner	Response within SLA; escalates on low-confidence or out-of-frontier signals.	Any resolution outside a pre-cleared script; a live sample audited weekly, not just escalations.
Monitoring tripwires	Commercial lead + FP&A lead	Alert within 15 minutes of breach.	Auto-hold release; covenant response.

The design principle embedded in every row: the agent owns the cycle time; the human owns the decision.

Inside the tool's capability frontier, users work markedly faster at higher quality; outside it, they do worse than unassisted peers — one controlled study measured a 19-point loss of correctness. The service-level definition is how you operationalise the frontier: specify exactly which outputs the agent ships autonomously (refresh, flag, draft, brief) and which require the human gate (publish, release, sign).

No workflow is unowned. An agent without a named owner drifts — thresholds stale, feeds break, flags fire into an empty room — and an unowned agent is worse than no agent, because it manufactures false confidence in monitoring that is not actually happening.

G.6 The Four Guardrails

Guardrail	Standard
	Anything customer-visible, partner-visible, financial, legal or scientific passes a human gate. Agents draft; humans sign. The named owner attests, in the decision journal, that they reviewed the output.
	Material claims in external content trace to verifiable sources. Every agent-generated external artifact carries a source manifest, and the signing human verifies at minimum every quantitative claim. A business that publishes an unsourced market claim at agent speed is manufacturing reputational risk at agent speed.

Guardrail	Standard
	No single AI vendor becomes a greater-than-25% single point of failure for critical workflows. An API deprecation, pricing shock or capability regression at your sole vendor is structurally identical to losing your largest account — a revenue-side event on the cost side.
	Agents hold credentials scoped to their service-level definition and nothing more; every agent action is logged against its owner's name; convenience never outranks the threat model. An agent with broad credentials and no logging is not leverage — it is an unauditable insider.

OPERATING DOCTRINE

The guardrail set must be teachable in one sentence, or it will not survive the first quarter of real operations. Here is the sentence: agents draft, humans sign, on anything that leaves the building, moves money, or touches a real person's live problem.

CONSEQUENCE FRAMING

One consumer-finance company's assistant was, by its own account, doing the work of roughly 700 agents — 2.3 million conversations a month, resolution time cut from eleven minutes to two. Fourteen months later the CEO was publicly stating the all-AI approach was not the right one and rehiring humans. The mechanism is the one the guardrails exist to catch: the system optimised for resolution speed and cost per contact, both easily measured, while whether the resolution was correct and whether the customer trusted it degraded quietly enough to take over a year to surface. Audit the cases the system was confident enough not to escalate — those are the ones that matter.

G.7 The Revenue-per-FTE Doctrine

Most firms treat revenue per FTE as an outcome they discover at year-end. Treat it as a design target, reviewed quarterly alongside gross margin, with the same seriousness.

Element	Standard
The doctrine	No additional W-2 headcount beyond a lean executive structure until the automation review fails.
The automation review	Every proposed hire must first fail a documented review: what agent workflow, vendor capability or process redesign was attempted, and what were its measured limits? The burden of proof sits on the headcount request, exactly as it sits on any capital request.
Why the burden sits there	Payroll is the most illiquid, highest-switching-cost vendor contract a firm ever signs. No other contract carries severance on exit, compounds at merit-cycle rates, embeds itself in the culture and resists termination in a downturn. A \$120,000 hire is not a \$120,000 decision; loaded and compounded it is a seven-figure decade-long commitment against a problem a \$2,000-a-month orchestration layer may solve.
The default order	Vendor and orchestration options before headcount — the evidence says bought, workflow-embedded solutions succeed roughly twice as often as internal builds.
The quarterly instrument	Three numbers on the same slide as gross margin: revenue per FTE against target; agent-automation coverage (share of recurring workflows with a live service-level definition); and the count of automation reviews that failed into a hire — each carrying a one-paragraph postmortem naming the measured limit that forced it.
Why the paper trail matters	A doctrine without a paper trail decays into a hiring freeze that breaks at the first urgent requisition. A

Element	Standard
	doctrine with a paper trail improves, because every failure teaches the organisation where the current automation frontier actually sits.
What this is not	It is not an argument for starving the business of people. It is an argument for spending people where people are the only answer — taste, judgment, relationships, irreversible calls — and refusing to spend them on refresh, reconcile, flag, draft.

PART H

The Velocity Dashboard

One dashboard, three panels, every metric traced to the driver tree, leading indicators outnumbering lagging ones two to one. Anything not on this page is either automated background monitoring or deleted. The dashboard's job is not to inform; it is to allocate attention — the scarcest resource in the operating system.

OPERATING DOCTRINE

The test for every tile is mechanical: if this moves today, does the P&L move within 90 days? If yes, it belongs on the weekly panels. If the answer is "the P&L tells me this already happened," it moves to monthly attestation or off the page entirely. A dashboard that is mostly outcomes is a history book with good typography.

A dashboard where everything is green is either a lie or a set of targets chosen to be unfailable. Both are management defects. Expect three to five tiles behind target at any time; that is the instrument working.

H.1 Panel One — Survivorship

Reviewed weekly in one minute of silence; attested monthly by name. The characteristic failure of survivorship metrics is not absence but inattention: everyone saw the number drift, and nobody was accountable for saying the sentence out loud.

Metric	Target	Cadence	Linked principle / P&L line
Months of runway (13-week rolling cash)	> 12 months	Weekly review / monthly attestation	Cash; survivorship precedes returns
Customer or channel concentration (largest exposure)	< 25% of net revenue; flag at 20%	Monthly attestation	Cash; uncapped downside
Position sizing vs. fractional Kelly	All bets $\leq$ cap	Quarterly audit	Enterprise value; Kelly discipline
Average discount depth	Below escalation tripwire	Weekly	Gross margin; hidden economics
Returns & deductions as % of gross	Within benchmark band; no single item above the item cap	Weekly tripwire	Gross margin; gross-to-net waterfall
Early-cohort cancels (months 1–3)	Flat or falling vs. prior cohorts	Weekly	Revenue; cohort economics
Cohort LTV (gross-margin basis, by vintage)	Rising by cohort; LTV:CAC $\geq$ 3:1	Monthly	S&M; LTV:CAC floor
Discretionary spend as % of gross revenue	Within band, not trending up	Monthly	S&M; hidden economics
Compliance / quality incidents	Zero open severity-1	Weekly tripwire	Enterprise value; reputation assets
Vendor / model concentration	No single provider > 25% single point of failure	Quarterly	Opex; survivorship covenant

H.2 Panel Two — Throughput (The War Room Scorecard)

The highest density of leading indicators on the dashboard. Design rule: each tile must be actionable in the meeting that reviews it. If the War Room cannot change a metric, the metric does not belong in the War Room.

Metric	Target	Cadence	Linked P&L line
First-purchase activation rate	Rising	Weekly	Revenue → active customers
30-day repeat-purchase rate, by cohort	Rising	Weekly	Revenue → purchase frequency
Subscription penetration	Within target band	Weekly	Revenue → frequency and LTV
Early-cohort cancels	Falling vs. prior cohort	Weekly	Revenue → churn
Reactivation / win-back rate	8–15% of recently lapsed	Weekly	Revenue at ~½ CAC
Plays shipped and adopted (one per week)	1/week; ≥ 70–80% adoption	Weekly	Revenue → execution quality
True promotion / spend ROI (60–90 day net contribution incl. decay)	Positive on every gated event	Per event, day-60 gate	Gross margin
Velocity per unit of distribution in priority accounts	≥ buyer hurdle; rising per unit	Weekly	Revenue → velocity
Weighted distribution	Expand only after velocity proves out	Monthly	Revenue → distribution

If you move only one metric on this dashboard, move the 30-day repeat rate. It is the single best leading indicator of durable lifetime value in any recurring-purchase business, and it is only visible per cohort. Note also the conditional target on distribution: distribution bought before velocity proves out is the most reliable way to purchase a write-off, which is why that row carries a condition rather than a growth number.

H.3 Panel Three — Velocity & Position

This panel measures the machine rather than the market: how fast the experimentation engine runs, and whether the firm's structural position is compounding. These are the metrics almost no company can currently answer — which is precisely why they belong on the page.

Metric	Target	Cadence	Linked principle
Experiments shipped per week	≥ 3	Weekly Growth Review	Learning velocity is the master metric
Cycle time per test (idea → verdict)	< 14 days	Weekly	Cadence is the delivery mechanism
Cost per test	Falling 2–5× vs. baseline	Monthly	AI as electricity
Win rate vs. one-third base rate	≈ 33%; kill rate ≥ 50%	Weekly verdicts	Data over narrative

Metric	Target	Cadence	Linked principle
Convexity ratio of pipeline	> 60% structurally unbounded	Quarterly Reset	Distribution realism
Inbound share of qualified pipeline	Rising	MBR	Preferential attachment
Structural-hole / new-network contacts	≥ 5 per week	Weekly	Exploit structural holes
Reputation assets shipped (case studies, talks, data releases)	≥ 4 per month	MBR	Compound reputation
Revenue per FTE	Rising vs. designed target	MBR	AI operating leverage
Threshold progress (documented wins toward the 2–3 threshold)	On pace	Quarterly Reset	Position for preferential attachment

None of these metrics print in this quarter's P&L. All of them determine what the P&L can print two years from now. Confirmation flows one way: the P&L never manages the War Room. The moment lagging metrics start dictating weekly action, the operating system has been inverted and the team is steering by the wake.

CONSEQUENCE FRAMING

Run without leading indicators and the sequence is always the same. Quarter one: revenue holds, because forward-buying and discount depth can hold any revenue line for a quarter. Quarter two: the repeat rate decays and spend climbs to defend the number. Quarter three: gross margin confirms what the cohorts said in quarter one — except now the fix costs a reset, a write-down and a leadership change. The largest players in consumer goods had dashboards through the 2021–2024 pricing cycle; after more than 20% cumulative price increases they grew revenue 1.2% with EBIT margins near a ten-year low. They had dashboards. They did not have instrumentation. The difference is 90 days.

PART I

Installation

The system installs the way it operates: gated and on the clock. Ninety days, four gates, no gate skipped. Every failed transformation shares one structural defect — the work was sequenced as a project (a strategy phase, a design phase, a rollout phase) rather than as an operating loop with evidence gates. Projects produce deliverables; loops produce installed behaviour. The distinction shows up eighteen months later, when the project binder is on a shelf and the loop is still running the Monday meeting.

Three design rules

- The clock is the constraint. Every phase has a hard gate date, and a gate review that produces "we need two more weeks" is a No-Go, not a deferral. Deferral is how 90-day plans become 90-week plans.
- Each gate converts activity into a capital-allocation decision. Scale dollars release only on evidence.
- The CEO's calendar is part of the installation: 8–10 hours per week in days 1–30, tapering to 4–6 hours by days 61–90. The taper is itself the test. If the CEO's time cannot come down, the system is not installing — it is being held up by hand.

I.1 The Four Phases

Window	Work	Gate (measurable)	Failure mode → countermeasure
	Stand up the war room, the weekly scorecard and the 7-day ship list. Publish RACI and the stop-doing list. Build the driver tree from the actual P&L and decompose until the binding constraint is visible in cohort data. Run the trust and confidence scan. Install survivorship instrumentation: 13-week cash forecast, concentration audit, gross-to-net waterfall, spend gate with pre-agreed rejection criteria.	Binding constraint named on one page with cohort evidence; scorecard v1 live with ≤ 7 metrics, each with owner and baseline; cash bleed quantified and at least one bleed-stopping action executed; RACI published.	The diagnostic expands to fill available curiosity — week three arrives with a beautiful segmentation and no decision. → The ten-day clock; Gate 1 convenes on day 14 regardless.
	Launch the first experiment slate against the constraint — two experiments, not twelve, every criterion pre-registered before launch. Stand up the automation layer for scorecard production and cohort decomposition. Baseline cost-per-test and revenue-per-FTE. Run the hiring scorecard and vendor matrix on any open decisions; freeze all procurement and offers lacking matrix scores.	Pre-registered movement in the target leading metric at or beyond the registered MDE against control; spend-gate rejection rate and spend-as-%-of-gross trending to plan; scale plan drafted with ring sequencing.	Criteria drift — the pilot misses its bar and the room renegotiates success after the fact. → The signed experiment doc is read aloud at the gate. A miss is a kill or a re-pilot, never a rename.
	Scale validated pilots ring by ring, killing failures publicly and cheaply. Where the constraint is behavioural, deploy a time-boxed incentive overlay that pays on scorecard metrics rather than gross volume — correcting the structural mismatch where the front line is paid on volume while finance manages margin. Install the performance cadence and run the first keeper-test review.	Unit economics hold at each ring; overlay ROI $\geq 1.5\times$ on attributable margin or it is redesigned or terminated; governance adopted — the weekly scorecard meeting has run four consecutive weeks with decisions logged.	Premature scale — ring one launched on enthusiasm, and a bad result in a major account costs trust that took eight weeks to rebuild. → Ring gating; the CEO personally delivers the "not yet."
	Codify every validated play into standard work.	The operating model runs two	Founder gravity — the CEO keeps

Window	Work	Gate (measurable)	Failure mode → countermeasure
	Sunset or formalise all temporary structures. Run the first Quarterly Reset: convexity audit, threshold-progress review, capital reallocation toward what the gates validated. Schedule the annual offsite; commit next quarter's experiment slate and spend calendar; publish the continuous-improvement cadence.	consecutive full cycles — scorecard meeting, ship list, gate review — without founder intervention; playbooks documented, trained and in use by people who were not in the pilots; gains sustained across both cycles within tolerance.	re-entering decisions the system is designed to make. → The taper is contractual; the CEO's gate role shifts from deciding to auditing decision quality, and any override is logged and reviewed at the next reset.

OPERATING DOCTRINE

Ninety days is not arbitrary. It is the longest an organisation will hold intensity without a visible win, and the shortest window in which four evidence gates can each be honoured. Treat a failed gate as the plan working, not the plan failing — the loop-backs are where the money is saved. The only unrecoverable failure is a gate that passes on narrative.

CONSEQUENCE FRAMING

Longitudinal work on cost-transformation programs finds 40–60% of implementations fail and 40–50% of the value erodes within two years without governance. The absence of a lock-in gate, not the absence of analysis, is what kills them. Gate 4 is the one most often skipped and the only one that determines whether any of the preceding work survives contact with next year.

I.2 First Monday Morning

Ten actions, in order, before the week gets away from you. None requires a consultant, a new system, or permission.

#	Action	Done when
1	Stand up the war room.	One physical or virtual room, one wall of scorecards, one owner. Decisions get made there or they do not get made.
2	Publish the single-threaded owner list.	Every initiative, every metric, every gate has exactly one name against it.
3	Build the driver tree from the actual P&L.	One tree, no orphan metrics, no residual "other" bucket.
4	Stand up the 13-week cash forecast.	Weekly, rolling, non-delegable.
5	Publish the stop-doing list.	Tail products, zombie programs, meetings with no decision output. Stopping is the fastest capacity you will ever create.
6	Name the binding constraint.	One constraint confirmed with cohort data — not five priorities confirmed with enthusiasm.
7	Write the first experiment slate against that constraint.	Pre-registered success and kill criteria signed before anything launches.
8	Institute the daily huddle and the 7-day ship list.	Fifteen minutes, standing, one scorecard. Cadence starts this week, not after the offsite.
9	Run the keeper test on the top twenty people and	Both directions, in writing. The "maybes" become development or

#	Action	Done when
	the top ten vendors.	exit plans within thirty days.
10	Ship the first structural-hole contact and the first reputation asset before noon.	One bridge-building outreach, one piece of authority-building content. The network and the reputation compound from day one or not at all.

Day one is not about being ready. It is about making the first rotation of the loop. The system compounds — but only from motion.

APPENDIX

Tear-Outs and Registers

Print these pages. They are designed to be used away from the document.

X.1 The Threshold Register

Every operating number in the system, in one place. Where a number is a benchmark band rather than a covenant, it is marked as such. Calibrate the bands to your industry at the Annual Offsite; never mid-quarter.

Domain	Threshold	Type
Cash runway	> 12 months on 13-week rolling forecast	Covenant
Single customer / channel concentration	Flag at 20%, breach at 25% of net revenue	Covenant
Single position vs. enterprise value	≤ 25%	Covenant
Vendor / model single point of failure	≤ 25% of critical workflows	Covenant
Open severity-1 compliance or quality items	Zero	Covenant
LTV:CAC by cohort	≥ 3:1	Floor
Overlay / incentive ROI at scale gate	≥ 1.5× on attributable margin	Gate
True promotion ROI window	60–90 days, net of post-event decay	Method
Discretionary trade / promotional spend	15–25% of gross revenue	Benchmark band
Partner deductions	5–15% of gross	Benchmark band
Returns and allowances	3–10% of gross	Benchmark band
Net revenue as % of gross	Typically 60–70%	Reality check
Post-event velocity, week one	40–60% of baseline; recovery by ~week four	Reality check
Experiment win rate	≈ 33%	Base rate
Experiment kill rate	≥ 50%	Floor
Tests shipped per week	≥ 3	Target
Cycle time per test	< 14 days	Target
Convexity ratio of active pipeline	> 60% structurally unbounded	Target
Convexity score for outside-zone entry	> 3.0 composite	Gate
Play adoption before redesign	≥ 70–80% within two weeks	Gate
WIP limit on active initiatives	5, with a displacement rule	Covenant

Domain	Threshold	Type
Candidate matrix composite	≥ 4.0 to hire; no exceptions	Gate
Vendor matrix composite	≥ 4.0 to pilot; 3.0–3.9 with written mitigations; < 3.0 pass; any single 1 is a veto	Gate
Vendor pilot before annual contract	30–60 days, paid, pre-registered metrics	Covenant
Mis-hire cost assumption (senior)	5–15× annual compensation	Planning assumption
Keeper test	Quarterly, in writing; two consecutive "no"s = execute within two weeks	Covenant
Reactivation economics	8–15% of recently lapsed recovered at ~½ blended CAC	Benchmark
Decision latency — Type 1 material	≤ 10 business days	SLA
Decision latency — Type 2 delegated	≤ 3 business days	SLA
Gate 4 definition of installed	Two consecutive clean cycles without founder intervention	Gate

X.2 The Failure-Mode Index

Symptom → diagnosis → the rule you are violating. Use this when something is wrong and you do not yet know what.

Symptom you observe	Actual diagnosis	Go to
Revenue is growing but margin is falling	Volume is being purchased with discount depth. The baseline is eroding underneath the top line.	D.4 tripwires; D.5 hidden economics; the standing question
Distribution or account count is up and revenue is flat	Distribution was bought before velocity was earned. The classic write-off purchase.	B.2 Iterate card; D.3; Panel Two conditional target
Everything on the dashboard is green	Targets were chosen to be unachievable, or the panel is mostly lagging indicators.	Part H opening; the 90-day mechanical test
The team says there are too many meetings	Decisions are being made at the wrong layer, so every forum re-litigates the layer below it.	C.1 control-system reading; C.2 meeting physics
Nothing gets approved and everything gets proposed	Approval gridlock and initiative sprawl are the same disease at different stages.	C.4 decision rights; WIP limit with displacement rule
Pilots keep getting extended	The pilot culture is a launch process with extra paperwork.	A.8 gate script; E.3 pre-registration; three-verdict rule
Growth review scales almost everything it tests	The pipeline is funding narrative, not discovering wins. Kill rate is below 50%.	E.5 velocity accounting; E.1 kill branch
A senior hire failed inside a year	Selection ran on conversation. Process yielded to persuasion.	F.3 matrix; F.1 work sample; separation post-mortem
Two people believe they own the same relationship	Shared ownership is unowned. The account gets attention when both are free, which is never.	C.2 single-threaded ownership

Symptom you observe	Actual diagnosis	Go to
A platform contract is delivering nothing but still being paid	No paid pilot, no pre-registered metrics, and switching cost was never priced.	G.1; G.2 auto-flags; G.3 thresholds
AI spend is rising with no P&L movement	Tools bolted onto unchanged workflows. The org chart was never redrawn.	G.4 orchestrate; G.5 agent org chart
An agent flag fired and nobody acted	The workflow is unowned. Unowned agents manufacture false confidence in monitoring that is not happening.	G.5; G.6 guardrails
The transformation gains faded after the leader stepped back	Phase 5 never happened. Nothing was codified into standard work.	B.2 Deploy card; I.1 Gate 4
Savings were delivered and the business still stalled	Cadence without a growth feedback loop is extraction with better hygiene.	C.3 MBR reinvestment test; A.3 zone router
A single account or vendor now decides your quarter	Concentration drifted past the covenant while everyone watched.	A.5 covenants; D.4 concentration tripwire

X.3 Standing Forms

Pre-mortem (one page, filed before signature)

- It is twelve months from now and this decision failed completely. Write the story in past tense.
- List the three most likely causes, ranked. Assign each a named owner.
- For each cause, name the leading metric that would surface it earliest, and the date by which it would be visible.
- State what you would need to see to abandon the decision, and who has authority to call it.
- Sign and date. File in the decision journal alongside the prediction.

Decision journal entry

- Decision, date, class (Type 1 material / Type 1 structural / Type 2).
- Predicted outcome with a number and a date. This is the field that makes the journal useful.
- Evidence relied on, and the base-rate check: what happened the last ten times a company like ours tried this.
- Matrix scores, if a matrix applied. Pre-mortem reference. Red-team memo reference.
- The name that signed. Review date.

Weekly ship list

- One line per commitment: owner, deliverable, binary done-or-not by the next war room.
- No percentages, no status colours, no partial credit.
- Last week's list is reviewed first, out loud, before this week's is written.

Gate memo (one page, every gate)

- The criterion, quoted verbatim from the signed pre-registration, with its date.
- The measured result against the required result.

- The verdict: kill, iterate or scale.
- The answer to "what would have to be true for us to override our own criterion?" — or the sentence "no override case was offered."
- The next commitment: owner, resource, date.

X.4 The Operating Identity

Run the inside like Toyota. Position the outside like Sequoia. Power everything like it is 1910 and you have just discovered the fractional-horsepower motor. Survive long enough for the avalanche, and hold convex exposure when it arrives.

What this system removes is not discomfort. It is ambiguity. You will always know, on any Tuesday, whether the business is compounding or drifting, because the dashboard will tell you before the P&L does. You will always know whether a bet is a bet or a hope, because the gate will have asked. You will always know whether a person or a vendor has earned the next quarter of your trust, because you will have written the answer down the quarter before.

Execution compounds. Install the war room this week.